



AUTOMATED DATA INGESTION AND PIPELINE SYSTEM
FOR AI MODEL TRAINING

MS. KWINYARUT POUNGSANGTHANAKUL
MR. PARIT LEELASETAWONG

A PROJECT SUBMITTED IN PARTIAL FULFILLMENT
OF THE REQUIREMENTS FOR
THE DEGREE OF BACHELOR OF ENGINEERING (COMPUTER ENGINEERING)
FACULTY OF ENGINEERING
KING MONGKUT'S UNIVERSITY OF TECHNOLOGY THONBURI
2026

Automated Data Ingestion and Pipeline System
for AI Model Training

MS. Kwinyarut Pounsangthanakul
Mr. Parit Leelasetawong

A Project Submitted in Partial Fulfillment
of the Requirements for
the Degree of Bachelor of Engineering (Computer Engineering)
Faculty of Engineering
King Mongkut's University of Technology Thonburi
2026

Project Committee

..... (Dr. Aye Hninn Khine)	Project Advisor
..... (Dr. Piyanit Ua-areemitr)	Committee Member
..... (Mr. Naveed Sultan)	Committee Member

Copyright reserved

Project Title	Automated Data Ingestion and Pipeline System for AI Model Training
Credits	18
Member(s)	MS. Kwinyarut Pounsangthanakul Mr. Parit Leelasetawong
Project Advisor	Dr. Aye Hninn Khine
Program	Bachelor of Engineering
Field of Study	Computer Engineering
Department	Computer Engineering
Faculty	Engineering
Academic Year	2026

Abstract

The growing demand for AI at Siam.AI Cloud requires large, high-quality datasets that can be processed reliably across many different formats. However, current data preparation workflows are fragmented, largely manual and error-prone especially for Thai-language data. This project addresses these challenges by building an end-to-end, open-source, on-premise pipeline that takes raw data in various formats and converts it into clean, standardized outputs ready for training large language models.

The system uses a Data Lakehouse design with MinIO for scalable object storage, PostgreSQL for metadata and processing logs and Apache Airflow for automating the entire workflow. The pipeline handles ingestion, cleaning, normalization, deduplication and quality validation across various data types, with a primary focus on the NECTEC corpus, covering formats such as JSONL, Parquet, CSV, TXT, PDF documents and audio recordings. For PDFs, the pipeline employs an OCR model deployed in order to extract text from scanned and digitally-encoded Thai documents through a dual-pass strategy, while attempting fast text extraction first and falling back to OCR when encoding issues are detected. Moreover, the audio pipeline leverages an ASR-Whisper model to transcribe Thai speech, incorporating a multi-level quality checking system capable of validating output even in the absence of ground truth labels. Pipeline health, storage usage and data quality are monitored in real time through Prometheus and Grafana dashboards.

The resulting system reduces manual effort, ensures consistent data quality across all processing stages, and produces standardized JSONL outputs ready for AI model training. It establishes a reusable foundation for Thai LLM development and future data engineering work at Siam AI Cloud.

Keywords: Data Ingestion / Data Lakehouse / Apache Airflow / OCR / ASR / Thai NLP

หัวข้อปริญญานิพนธ์	ระบบการนำเข้าข้อมูลและสายงานประมวลผลแบบอัตโนมัติ สำหรับการฝึกแบบจำลองปัญญาประดิษฐ์
หน่วยกิต	18
ผู้เขียน	นางสาว กวินณรัตน์ พ่วงแสงธนกุล นาย ปรีดี ลีลาเศรษฐวงศ์
อาจารย์ที่ปรึกษา	ดร. เอ หนิน ชื่น
หลักสูตร	วิศวกรรมศาสตรบัณฑิต
สาขาวิชา	วิศวกรรมคอมพิวเตอร์
ภาควิชา	วิศวกรรมคอมพิวเตอร์
คณะ	วิศวกรรมศาสตร์
ปีการศึกษา	2569

บทคัดย่อ

การเติบโตอย่างรวดเร็วของเทคโนโลยีปัญญาประดิษฐ์ใน Siam.AI Cloud ส่งผลให้เกิดความต้องการชุดข้อมูลขนาดใหญ่และมีคุณภาพสูง ซึ่งสามารถประมวลผลข้ามรูปแบบที่หลากหลายได้อย่างมีประสิทธิภาพ อย่างไรก็ตาม กระบวนการเตรียมข้อมูลในปัจจุบันยังมีลักษณะกระจัดกระจาย ต้องอาศัยการดำเนินการแบบแมนนวลและมีความเสี่ยงต่อความผิดพลาดสูงโดยเฉพาะอย่างยิ่งกับข้อมูลภาษาไทย โครงการนี้มีวัตถุประสงค์เพื่อแก้ไขปัญหาดังกล่าวผ่านการพัฒนาระบบสายงานประมวลผลแบบครบวงจร (End-to-end Pipeline) ประเภทโอเพนซอร์ส (Open-source) ที่ติดตั้งใช้งานภายในพื้นที่ปฏิบัติการ เพื่อรวบรวม แปลง และเตรียมข้อมูลดิบจากหลากหลายรูปแบบให้เป็นผลลัพธ์มาตรฐานที่สะอาดและพร้อมใช้งานสำหรับการฝึกแบบจำลองภาษาขนาดใหญ่ (Large Language Models: LLMs)

ระบบถูกออกแบบภายใต้แนวคิดคลังข้อมูลแบบผสมผสานระหว่างแหล่งเก็บข้อมูลขนาดใหญ่และคลังข้อมูลเชิงวิเคราะห์ (Data Lake-house) โดยใช้ MinIO เป็นระบบจัดเก็บข้อมูลแบบวัตถุที่รองรับการขยายตัว ใช้ฐานข้อมูลเชิงสัมพันธ์ PostgreSQL สำหรับจัดการเมตาดาตาและบันทึกการประมวลผล และใช้ Apache Airflow ในการควบคุมการทำงานแบบอัตโนมัติทั้งระบบ รองรับขั้นตอนการนำเข้าข้อมูล การทำความสะอาดข้อมูล การปรับมาตรฐาน การจัดข้อมูลซ้ำ และการตรวจสอบคุณภาพ ครอบคลุมข้อมูลหลายประเภท โดยมุ่งเน้นที่คลังข้อมูล NECTEC ซึ่งประกอบด้วยรูปแบบไฟล์ JSONL, Parquet, CSV, TXT, เอกสาร PDF และไฟล์เสียง

สำหรับการประมวลผลเอกสาร PDF ระบบได้ประยุกต์ใช้แบบจำลองการรู้จำอักขระด้วยแสง (OCR) เพื่อสกัดข้อความจากเอกสารภาษาไทยทั้งแบบภาพสแกนและแบบเข้ารหัสดิจิทัลผ่านกลยุทธ์การประมวลผลแบบสองขั้นตอน (Dual-pass Strategy) นอกจากนี้สายงานประมวลผลข้อมูลเสียงได้ใช้แบบจำลองรู้จำเสียงพูดอัตโนมัติ (ASR) Whisper ในการถอดความเสียงพูดภาษาไทย ร่วมกับระบบตรวจสอบคุณภาพแบบหลายระดับ ทั้งนี้สถานะการทำงานของระบบ ปริมาณการใช้พื้นที่จัดเก็บ และคุณภาพข้อมูลจะได้รับการเฝ้าระวังแบบเวลาจริงผ่านระบบแดชบอร์ด Prometheus และ Grafana ผลลัพธ์ของระบบช่วยลดภาระการทำงานแบบแมนนวล เพิ่มความสม่ำเสมอของคุณภาพข้อมูลในทุกขั้นตอน และเป็นโครงสร้างพื้นฐานที่สามารถนำไปประยุกต์ใช้สำหรับการพัฒนาแบบจำลองภาษาไทยและงานด้านวิศวกรรมข้อมูลของ Siam.AI Cloud ในอนาคต

คำสำคัญ: การนำเข้าข้อมูล / คลังข้อมูลแบบผสมผสาน / Apache Airflow / OCR / ASR / การประมวลผลภาษาไทย

ACKNOWLEDGMENTS

The authors wish to express their profound gratitude to project advisor Dr. Aye Hninn Khine for her invaluable guidance, scholarly insight and steadfast support throughout the duration of this project. Her expertise, constructive feedback and unwavering commitment to academic excellence have been instrumental in shaping the quality, direction and successful completion of this work.

We extend our sincere appreciation to Siam.AI Cloud Co., Ltd. for providing the opportunity to undertake this project under the Work-Integrated Learning (WIL) program. The access to enterprise-grade infrastructure and real-world data engineering environments has significantly enriched the practical relevance and technical depth of this study. Our deepest thanks are also conveyed to our mentor Mr. Abhibhu Tachaapornchai whose exceptional mentorship, technical guidance and continuous encouragement played a critical role in the development and refinement of the proposed system.

The authors further wish to acknowledge the project committee members for their thoughtful recommendations and rigorous evaluations which greatly contributed to the improvement and academic rigor of this work. Finally, we express our heartfelt gratitude to our families, friends and colleagues for their encouragement, understanding and unwavering support throughout the course of this project.

CONTENTS

	PAGE
ABSTRACT	ii
THAI ABSTRACT	iii
ACKNOWLEDGMENTS	iv
CONTENTS	v
LIST OF TABLES	viii
LIST OF FIGURES	ix
LIST OF TECHNICAL VOCABULARY AND ABBREVIATIONS	x
CHAPTER	
1. INTRODUCTION	1
1.1 Project Introduction and Problem Statement	1
1.2 Objectives	1
1.3 Scope and Limitations	2
1.3.1 Data Types and Ingestion	2
1.3.2 Architecture and Tools	2
1.3.3 Deliverables	2
1.3.4 Phase 1: NECTEC Dataset Integration (Primary Focus)	2
1.3.5 Phase 2: Audio Data Pipeline and System Enhancement	3
1.3.6 Limitations on Data Types (Out of Scope)	3
1.3.7 Limitations Related to Target System	3
1.3.8 Out-of-Scope Items	4
1.4 Project Plan	4
1.4.1 August (Preparation and Research)	4
1.4.2 September (Environment Setup and Baseline Ingestion)	4
1.4.3 October (Automation and PDF Pipeline Development)	4
1.4.4 November (Integration and Monitoring Setup)	4
1.4.5 December (Structured and Semi-Structured Expansion)	4
1.4.6 January (Unstructured Data Expansion)	5
1.4.7 February (Audio Data Pipeline Phase 1)	5
1.4.8 March (Audio Data Pipeline Phase 2)	5
1.4.9 April (Advanced Monitoring and System Refinement)	5
1.4.10 May (Final Refinement and Documentation)	5
1.5 Gantt Chart	5
1.6 Deliverables	6
1.6.1 First Semester (August–December)	6
1.6.2 Second Semester (January–May)	6
2. BACKGROUND, THEORY AND RELATED RESEARCH	7
2.1 Background	7
2.2 Development Tools	7
2.2.1 MinIO	7
2.2.2 Apache Airflow	7
2.2.3 PyThaiNLP	8
2.2.4 OCR Models (Tesseract, Typhoon, EasyOCR, PaddleOCR)	8
2.2.5 PostgreSQL	8
2.2.6 Monitoring Tools (Airflow UI, Prometheus, Grafana)	8
2.2.7 ASR Model (Typhoon Isan ASR Whisper)	9

2.2.8	Visual Studio Code	9
2.2.9	FortiClient VPN	9
2.2.10	WangchanBERTa-finetuned-sentiment	9
2.3	Existing Solutions and Products in the Market	9
2.3.1	Cloud Provider Data Platforms (AWS, GCP, Azure)	9
2.3.2	Open-Source Data Ingestion and Integration Tools (NiFi, Airbyte)	10
2.3.3	Positioning of the Proposed System	10
2.4	Preliminary Design Based on Backgrounds	11
3.	METHODOLOGY AND DESIGNS	15
3.1	Introduction	15
3.2	Functional Requirements	15
3.2.1	Data Ingestion Requirements	15
3.2.2	Data Cleaning and Transformation Requirements	15
3.2.3	Output and Storage Requirements	15
3.2.4	Orchestration and Monitoring Requirements	15
3.3	System Design	16
3.3.1	Architectural Overview	16
3.3.2	Containerized Service Architecture	16
3.4	Data Lake Structure	16
3.4.1	MinIO Structure	17
3.4.2	Bucket Versioning and Replication	17
3.5	Data Warehouse	17
3.6	Orchestration	20
3.6.1	JSONL Pipeline	20
3.6.2	Unstructured PDF OCR Pipeline	22
3.6.3	Unstructured Audio ASR Pipeline	25
3.7	Monitoring	28
3.7.1	Metrics Collection with Prometheus	28
3.7.2	Grafana Dashboard	29
4.	RESULTS AND DISCUSSION	30
4.1	Data Lake	30
4.2	Unstructured Pipeline	31
4.2.1	PDF Pipeline	31
4.2.2	Audio Pipeline	33
4.3	JSONL Pipeline	36
4.3.1	Main Pipeline Execution	36
4.3.2	Sentiment Processing Pipeline	37
4.4	Database	38
4.4.1	Datasets	38
4.4.2	Documents	39
4.4.3	Labels	40
4.4.4	preprocessing_logs	41
4.5	Dashboard	41
4.5.1	Airflow Dashboards	42
4.5.2	MinIO Dashboards	42
4.5.3	PostgreSQL	43
4.5.4	Data Quality	43
5.	CONCLUSION	44
5.1	Conclusion	44

5.2	Problems, Obstacles and Solutions	45
5.2.1	Inconsistent Data Format Across Sources	45
5.2.2	Data Quality Issues in Thai Language Processing	45
5.2.3	System Performance and Scalability Issues	46
5.3	Future Work and Suggestions	46
REFERENCES		47

LIST OF TABLES

TABLE	PAGE
1.1 Gantt chart — First semester (August–December)	5
1.2 Gantt chart — Second semester (January–May)	6
2.1 Comparison of existing data platforms vs. the proposed Siam.AI pipeline	10
3.1 <code>datasets</code> table schema	18
3.2 <code>documents</code> table schema	18
3.3 <code>conversations</code> table schema	18
3.4 <code>messages</code> table schema	19
3.5 <code>labels</code> table schema	19
3.6 <code>preprocessing_logs</code> table schema	19
3.7 <code>evaluation_metrics</code> table schema	19
4.1 Overall unstructured pipeline output	31
4.2 PDF dataset overview	31
4.3 PDF document-level results	31
4.4 PDF page-level results	32
4.5 PDF extraction method results	32
4.6 Audio dataset overview	33
4.7 Audio acceptance results	33
4.8 Audio output volume	33
4.9 CER distribution of accepted records	33
4.10 CER distribution of filtered records	34
4.11 Audio duration profile of accepted records	34
4.12 Per-shard result consistency	34
4.13 Audio format compatibility test results	34
4.14 Format compatibility replication results	35
4.15 Combined format-coverage summary	35
5.1 Project completion status	45

LIST OF FIGURES

FIGURE	PAGE
2.1 Automated Data Ingestion and Pipeline System Architecture	11
2.2 NVIDIA B200 GPUs	11
2.3 MinIO Object Storage (Data Lake Interface)	12
2.4 pgAdmin 4 Interface for PostgreSQL Data Warehouse	13
2.5 Apache Airflow Interface for Data Pipeline Connection	13
3.1 MinIO Interface	16
3.2 Database schema for PostgreSQL	17
3.3 Apache Airflow graph view DAGs	20
3.4 Apache Airflow Build environment task detail	21
3.5 Apache Airflow Sensor_demo DAG detail	21
3.6 Apache Airflow graph view of pdf_ocr_pipeline	22
3.7 Apache Airflow process_pdf task detail	23
3.8 Apache Airflow pdf_ocr_pipeline DAG run detail	24
3.9 Apache Airflow graph view of audio_asr_pipeline	25
3.10 Apache Airflow process_audio_file task detail	26
3.11 Apache Airflow audio_asr_pipeline DAG run detail	27
3.12 Prometheus metric connector	28
3.13 Grafana Dashboard	29
4.1 Datalake Layer	30
4.2 NECTEC Datasets in MinIO	30
4.3 Main pipeline DAG	36
4.4 Sentiment Processing Pipeline	37
4.5 Database schema in PostgreSQL	38
4.6 Dataset tables	39
4.7 Document Table	39
4.8 Document Table (detail view)	40
4.9 Labels Table	40
4.10 Preprocessing_logs Table	41
4.11 Airflow Dashboards	42
4.12 MinIO Dashboards	42
4.13 PostgreSQL Dashboards	43
4.14 Data Quality Dashboards	43

LIST OF TECHNICAL VOCABULARY AND ABBREVIATIONS

AI	=	Artificial Intelligence
ASR	=	Automatic Speech Recognition
CER	=	Character Error Rate
DAG	=	Directed Acyclic Graph
ETL	=	Extract, Transform, Load
GPU	=	Graphics Processing Unit
LLM	=	Large Language Model
NECTEC	=	National Electronics and Computer Technology Center
NLP	=	Natural Language Processing
OCR	=	Optical Character Recognition
WER	=	Word Error Rate
WIL	=	Work-Integrated Learning

CHAPTER 1 INTRODUCTION

1.1 Project Introduction and Problem Statement

Siam.AI Cloud is a technology company specializing in artificial intelligence (AI) and cloud-native solutions which focus on enabling organizations to harness data-driven innovation. The company also provides infrastructure and services for large-scale AI model development, including GPU-powered compute clusters (e.g., NVIDIA Blackwell B200/H100/H200), scalable object storage and open-source integration. With expertise in machine learning research and enterprise deployment, Siam.AI Cloud plays a vital role in advancing AI adoption across industries in Southeast Asia.

Against this backdrop, Siam.AI Cloud recognizes that the rapid expansion of AI requires large, high-quality datasets to train and fine-tune models. However, managing diverse types of data remains a major challenge. Preparing AI-ready datasets without automation is not only slow but also highly resource-intensive as engineers must repeatedly collect, clean and transform raw information from different sources using manual processes. This approach is error-prone which leads to issues such as missing values, inconsistent formatting, duplicated records or incomplete metadata tagging, all of which can degrade the quality of datasets and negatively impact downstream AI model performance. Additionally, as the volume, velocity and variety of data continue to grow, manual preparation methods cannot keep up. Tasks that may take days with small datasets could extend to weeks or months when scaled to enterprise or real-time environments. Without automation, organizations face bottlenecks that delay experimentation and slow the deployment of AI solutions. In contrast, an automated ingestion and preparation pipeline ensures faster processing, greater consistency and scalability, allowing raw data to be continuously transformed into standardized AI-ready formats that support efficient training and fine-tuning of machine learning models.

In order to solve this problem, this project proposes developing an Automated Data Ingestion and Pipeline System for AI Model Training. The pipeline will automate ingestion, cleaning, transformation and storage of multiple data types within a unified Data Lakehouse architecture. It will use MinIO for scalable object storage, PostgreSQL for metadata and analytics and Apache Airflow for workflow orchestration. Moreover, the transformation modules will handle tasks such as normalization, deduplication, metadata tagging and enrichment using tools. To ensure reliability, the system will include monitoring dashboards for pipeline health and performance tracking.

1.2 Objectives

- Develop an automated data ingestion and ETL pipeline that can process various types of data.
- Design a Data Lake using MinIO to store both raw and processed datasets.
- Use PostgreSQL as a Data Warehouse for querying, metadata management and analytical workloads.
- Implement data cleaning, normalization, deduplication and metadata tagging to improve dataset quality and usability.
- Produce standardized AI-ready data formats such as JSONL, TFRecords and Parquet optimized for machine learning pipelines.
- Provide a real-time monitoring dashboard to track pipeline health, system performance and dataset status.
- Ensure the overall system is reliable, scalable and suitable for large-scale AI training workflows.

1.3 Scope and Limitations

1.3.1 Data Types and Ingestion

- Structured: CSV, Excel, SQL (clean and load into Data Warehouse)
- Semi-structured: JSON, XML, YAML (normalize, flatten and convert into standardized formats)
- Unstructured: Text, PDFs, Images and Audio

1.3.2 Architecture and Tools

- Data Lake: MinIO (S3-compatible object storage)
- Data Warehouse: PostgreSQL (querying and metadata)
- Orchestration: Apache Airflow (DAGs for automation)
- Ingestion: Python scripts
- Cleaning/Transform: Pandas, PyArrow, OCR, ASR
- Monitoring: Prometheus, Grafana, Airflow UI

1.3.3 Deliverables

- Automated pipeline supporting all data types
- Data Lake with raw, cleaned and processed layers
- Metadata schema for datasets, documents and messages
- Repository of standardized AI-ready datasets (JSONL, Parquet, TFRecords)
- Monitoring dashboard for pipeline health and dataset tracking
- Documentation and final demo system

1.3.4 Phase 1: NECTEC Dataset Integration (Primary Focus)

- NECTEC dataset used as the initial and primary case study
- Handles diverse Thai corpora formats: JSONL, TXT, CSV, PDF
- Preprocessing tasks: ingestion into MinIO, cleaning, normalization, encoding fixes, duplicate removal with checksums, filtering with Thai-character ratios and export into standardized JSONL/Parquet
- Metadata stored in PostgreSQL (datasets, documents, preprocessing logs, validation status)
- Output: a consistent, high-quality, AI-ready Thai corpus for LLM training

In the first phase, the work focuses on integrating and preparing the NECTEC dataset as the initial case study. The NECTEC corpus consists mainly of Thai text data provided in different formats such as single JSONL corpora, domain-based folders containing index files, raw TXT documents, and file-to-URL mapping CSVs. Some sources also appear in paired CSV and JSONL formats. The processing in this phase includes ingestion of these files into MinIO, cleaning and normalization with a Python-based cleaner, removal of duplicates using checksums, filtering based on minimum Thai character ratios, and finally exporting standardized cleaned JSONL or Parquet files. Metadata such as dataset information, documents, preprocessing logs, and validation status will be stored in PostgreSQL according to the defined schema. The output of this phase is a consistent set of AI-ready Thai corpora that can be used for training Large Language Models (LLMs) or other AI models.

1.3.5 Phase 2: Audio Data Pipeline and System Enhancement

- Develop an audio data ingestion and preprocessing pipeline
- Apply Automatic Speech Recognition (ASR) for speech-to-text conversion
- Perform cleaning, normalization, and formatting of transcribed text
- Store audio metadata and transcription results in PostgreSQL
- Enhance and optimize the monitoring dashboard for better visualization and usability

In the second phase, the pipeline is extended to support audio datasets by developing ingestion workflows that handle audio file formats and integrate speech-to-text conversion using ASR models. The processing steps include audio ingestion into the data lake, conversion of speech into text using ASR, followed by text cleaning, normalization, and standardization into AI-ready formats such as JSONL or Parquet. Audio source, transcription logs, and processing status are stored in PostgreSQL alongside the existing schema. This phase also includes optimization of the monitoring dashboard to improve system observability.

1.3.6 Limitations on Data Types (Out of Scope)

The following are excluded from this project due to time, infrastructure or complexity constraints:

- Video datasets (require heavy GPU decoding and frame-level OCR/S2T)
- Streaming data (the company does not yet have a high volume of logs or event-based data that requires real-time processing)
- 3D / point-cloud data (not relevant to NECTEC and requires specialized libraries)
- Highly structured enterprise systems (SAP, ERP logs — no access to real data)
- Large-scale real-time IoT sensor streams (beyond current infrastructure)
- Multilingual OCR for non-Thai languages (scope focuses on Thai preprocessing)

The pipeline is optimized for Thai text-centric datasets, especially because NECTEC corpora are text-heavy and critical for Thai LLM training. Supporting additional complex modalities would exceed the academic timeframe and require specialized preprocessing not aligned with this project's objectives.

1.3.7 Limitations Related to Target System

Although the pipeline is designed to be extensible, the initial implementation is tailored specifically for the NECTEC dataset and Siam.AI's internal infrastructure:

- Schema dependence: preprocessing and metadata schemas are optimized for text datasets.
- OCR optimization for Thai (Typhoon OCR 7B): applying it to multilingual OCR workloads would require retraining or swapping OCR backends.
- ASR optimization for Thai (Typhoon Isan ASR Whisper): tuned for Thai/Isan, multilingual ASR would require swapping backends.
- NECTEC-oriented cleaning rules: Thai-character ratio filtering and domain-based normalization are specific to NECTEC's structure.
- Infrastructure-specific design: MinIO object storage, PostgreSQL metadata warehouse, Airflow orchestrator.

1.3.8 Out-of-Scope Items

- Direct AI/LLM training
- Model evaluation and fine-tuning
- Deployment of AI services
- Production-grade cluster scaling or multi-region redundancy
- Long-term storage governance (GDPR, data retention policies)

1.4 Project Plan

1.4.1 August (Preparation and Research)

The focus is on establishing a solid foundation for the project. The team will research and finalize the technology stack including MinIO, PostgreSQL, Apache Airflow and Python ETL libraries (Pandas, PyArrow). Monitoring tools such as Prometheus and Grafana will also be reviewed. The system architecture will be designed and the schema defined based on requirements for handling structured, semi-structured and unstructured data. A dataset catalog from NECTEC will be created by mapping available file types. Data quality rules will be specified for each file type.

1.4.2 September (Environment Setup and Baseline Ingestion)

September will be dedicated to building the project's baseline environment. The development environment will be set up with MinIO, PostgreSQL, Airflow and pgAdmin. The team will create initial ingestion and cleaning pipelines for structured (.csv) and semi-structured (.jsonl) data. In parallel, OCR models Tesseract, TyphoonOCR, PaddleOCR and EasyOCR will be compared for accuracy, speed and robustness with Thai datasets.

1.4.3 October (Automation and PDF Pipeline Development)

October will focus on bringing automation into the system. Ingestion and cleaning processes for structured and semi-structured data will be fully automated using Apache Airflow DAGs. The team will also develop the PDF ingestion and validation pipeline using the best OCR model identified in September. The PDF pipeline will be tested on NECTEC datasets that require OCR processing.

1.4.4 November (Integration and Monitoring Setup)

The PDF pipeline will be integrated into the existing Airflow DAGs, creating a complete end-to-end system. System monitoring will be established using Airflow's UI, complemented by Prometheus and Grafana. Log collection and error handling mechanisms will be added to ensure the system can detect and recover from common failures.

1.4.5 December (Structured and Semi-Structured Expansion)

With the core pipeline in place, December will expand ingestion to cover additional structured and semi-structured data sources such as CSV, Excel, SQL dumps, JSON, XML and YAML. Standardization of outputs will be enforced, converting all structured and semi-structured data into uniform formats such as Parquet or JSONL.

1.4.6 January (Unstructured Data Expansion)

January will focus on more complex unstructured data such as large text corpora, images and PDFs. The team will integrate OCR for PDFs and images, and NLP metadata tagging for text datasets. Outputs will be standardized into TXT, JSONL or TFRecord formats.

1.4.7 February (Audio Data Pipeline Phase 1)

The project will begin extending the pipeline to support audio data ingestion and preprocessing. The focus will be on designing the initial audio pipeline that can ingest audio files into MinIO and organize them into raw and processed storage zones. The team will also study and select a suitable ASR model.

1.4.8 March (Audio Data Pipeline Phase 2)

March will continue with the development of the audio data pipeline, focusing on integrating the selected ASR model into the existing workflow. Audio files will be processed through the ASR component to generate text transcriptions, which will then be cleaned, normalized and converted into standardized AI-ready formats such as JSONL or Parquet.

1.4.9 April (Advanced Monitoring and System Refinement)

April will focus on enhancing monitoring and reliability. Detailed dashboards will be built in Grafana to visualize ingestion rates, error counts and system performance. Prometheus alerts will be configured to notify the team of failures or abnormal activity.

1.4.10 May (Final Refinement and Documentation)

The final month will focus on stabilizing and documenting the system. Monitoring dashboards will be finalized with detailed visualizations of both system and data quality metrics. The pipelines will undergo performance testing and stress testing to confirm scalability. Comprehensive documentation will be prepared and a final demo of the system will be presented.

1.5 Gantt Chart

The high-level project timeline is summarized in Table 1.1 (first semester) and Table 1.2 (second semester).

Table 1.1 Gantt chart — First semester (August–December)

Task	Aug	Sep	Oct	Nov	Dec
Finalize technology stack and system architecture	X				
Database schema design and dataset catalog	X				
Define data quality rules	X				
Set up environment		X			
Implement database schema		X			
Baseline ingestion pipeline (.csv, .jsonl, .txt)		X			
OCR model evaluation		X	X		
Automate structured/semi-structured ingestion			X		
Develop PDF ingestion pipeline with OCR			X		
Test NECTEC PDF datasets			X	X	
Integrate PDF pipeline into full system				X	
Set up monitoring				X	
Implement logging and error handling				X	
Expand ingestion to Excel, SQL, JSON, XML, YAML					X
Standardize outputs					X
Extend validation rules					X

Table 1.2 Gantt chart — Second semester (January–May)

Task	Jan	Feb	Mar	Apr	May
OCR integration for images/PDF text	X				
NLP metadata tagging for text datasets	X				
AI-ready unstructured data outputs	X				
Set up audio data ingestion		X			
Develop audio preprocessing pipeline		X			
Integrate ASR model			X		
Store audio metadata and transcription results			X		
Apply validation rules for audio outputs			X		
Integrate audio pipeline with Airflow workflow			X		
Build detailed dashboards and refine error handling				X	
Optimize Airflow DAGs				X	
Finalize dashboards with data metrics					X
Performance and scalability test reports					X
Complete project documentation and Final Demo					X

1.6 Deliverables

1.6.1 First Semester (August–December)

- Complete ingestion, cleaning, validation and automation for NECTEC datasets (.csv, .jsonl, .txt, .pdf).
- Integrate OCR for Thai text and automate workflows with Airflow.
- Deliver an end-to-end pipeline storing processed NECTEC data in MinIO and PostgreSQL.

1.6.2 Second Semester (January–May)

- Extend the system to handle other data types.
- Integrate ASR for Thai audio and automate workflows with Airflow.
- Build monitoring and alerting dashboards (Grafana, Prometheus) with auto error handling.
- Finalize project with performance testing, documentation and full system demo.

CHAPTER 2 BACKGROUND, THEORY AND RELATED RESEARCH

2.1 Background

Artificial Intelligence (AI) and machine learning models need large amounts of good-quality data to work effectively. In practice, the data that organizations collect comes in many different forms such as spreadsheets, databases, JSON files, text documents, scanned PDFs, and even audio recordings. These different formats are often inconsistent and may contain errors or unnecessary information. Because of this, the data cannot be used directly for training AI models. In order to solve this problem, many organizations build data ingestion and preparation pipelines. These pipelines can collect data from different sources, clean and organize it, and store it in a standard format that can be used for AI training. At Siam.AI, the advantage is that the company already owns powerful computing resources, including GPU servers and large storage systems. This means the pipeline can be built completely with open-source tools without depending on external cloud providers. By using its own infrastructure, Siam.AI can reduce costs, keep data secure and have full control of the system. The goal of this project is to design and implement such a pipeline that can handle different types of data, apply cleaning steps and prepare data that is ready for AI model training in a reliable and repeatable way.

2.2 Development Tools

This project must run on the company's existing on-premise infrastructure and comply with their requirement to use open-source tools where possible. Therefore, the technology stack was initially suggested by the company. However, each component was evaluated against the project requirements and compared with alternatives. The selected tools were chosen because they best match these requirements while remaining compatible with the existing environment.

2.2.1 MinIO

MinIO is an open-source object storage system compatible with the Amazon S3 standard. It provides scalable storage for unstructured and semi-structured data and can handle both many small files and very large datasets efficiently. In this project, MinIO is used as the data lake, where raw and cleaned data from multiple sources are stored before and after preprocessing. The main project requirements for storage are on-premise deployment and no public cloud for privacy and cost reasons. Compared with alternatives such as AWS S3, Google Cloud Storage and HDFS, MinIO satisfies these constraints while still being lightweight and easy to deploy with Docker on the company server.

2.2.2 Apache Airflow

Apache Airflow is an open-source platform for authoring, scheduling and monitoring data workflows. Pipelines are represented as Directed Acyclic Graphs (DAGs), allowing explicit definition of task dependencies, retries and scheduling policies. In this project, Airflow automates the ingestion, cleaning and transformation processes, coordinating interactions between MinIO, Python preprocessing scripts and PostgreSQL. Simple solutions such as cron jobs offer limited support for complex dependencies and lack a central monitoring interface. Alternative workflow tools (Prefect, Dagster) are viable but are not currently part of the company's standard stack. Airflow is widely adopted in industry and provides a mature ecosystem of operators and monitoring features.

2.2.3 PyThaiNLP

PyThaiNLP is an open-source Python library developed for processing Thai text. It provides important functions such as tokenization, word and sentence segmentation, part-of-speech tagging, named entity recognition and text normalization. These tools are essential because Thai writing does not normally use spaces between words, making it difficult for computers to split text correctly without specialized methods. In this project, PyThaiNLP is applied to clean and preprocess Thai datasets before storing them for AI model training.

2.2.4 OCR Models (Tesseract, Typhoon, EasyOCR, PaddleOCR)

Optical Character Recognition (OCR) is an essential tool for extracting machine-readable text from scanned documents, PDFs and image files. In this project, OCR is used to transform unstructured visual data into structured text. Since different OCR frameworks offer varying levels of accuracy, speed and language support, this project evaluated and compared multiple models before final integration.

Tesseract is one of the most widely used open-source OCR engines, maintained by Google. It supports many languages including Thai and English and provides a stable, lightweight solution. Typhoon OCR is a model developed specifically with Thai language recognition in mind. It offers better segmentation for Thai characters compared to general OCR frameworks. EasyOCR is a deep learning-based OCR library that provides multilingual support and is easy to integrate with Python applications. PaddleOCR, developed by the PaddlePaddle community, is a high-performance OCR framework that supports over 80 languages and is known for its accuracy in complex layouts and noisy images.

In this project, these four OCR frameworks were benchmarked and compared using criteria such as accuracy on Thai and English text, processing speed, resource consumption and ease of integration with the pipeline. The best-performing model was then selected and deployed as part of the unstructured data processing workflow.

2.2.5 PostgreSQL

PostgreSQL is an open-source relational database management system that supports advanced querying, indexing and native JSON/JSONB data types. In this project, PostgreSQL serves as the data warehouse, storing dataset metadata, document records, preprocessing logs and data quality metrics. The warehouse must manage both structured information (identifiers, timestamps, counts) and semi-structured metadata. Compared with other relational databases such as MySQL, PostgreSQL offers stronger support for JSONB, extensible indexing strategies and analytical SQL features. Cloud-based warehouses (BigQuery, Snowflake) were not chosen due to the requirement for an on-premise, license-free solution.

2.2.6 Monitoring Tools (Airflow UI, Prometheus, Grafana)

Monitoring is essential for maintaining the reliability and performance of the pipeline. Airflow UI provides detailed information on DAG and task execution, including logs and retry status. Prometheus is employed to collect metrics from services such as PostgreSQL, Airflow and MinIO, while Grafana is used to visualize these metrics via dashboards. Together, these tools support real-time monitoring, troubleshooting and alerting. The combination of Prometheus and Grafana is widely adopted in industry, open source and already used by Siam.AI for other internal systems.

2.2.7 ASR Model (Typhoon Isan ASR Whisper)

Typhoon Isan ASR Whisper is an open-source Automatic Speech Recognition (ASR) model developed for transcribing Thai speech, with a specific focus on the Isan dialect. It is a fine-tuned model based on the Whisper architecture and is optimized for offline speech-to-text conversion, making it suitable for processing audio datasets without relying on third-party cloud services. The model improves transcription accuracy for regional Thai speech, which is often difficult for general ASR systems to recognize correctly. In this project, Typhoon Isan ASR Whisper is used as part of the audio data pipeline to convert raw audio files into text, which is then cleaned, normalized, validated and converted into standardized AI-ready formats such as JSONL or Parquet.

2.2.8 Visual Studio Code

Visual Studio Code (VS Code) is a lightweight, open-source integrated development environment (IDE) developed by Microsoft. It supports multiple programming languages, extensions, integrated terminal access and Git integration. In this project, VS Code is used as the primary editor for writing Python preprocessing scripts, designing Airflow DAGs and preparing SQL queries for PostgreSQL.

2.2.9 FortiClient VPN

FortiClient VPN is a security application developed by Fortinet that provides secure Virtual Private Network (VPN) connectivity. It enables users to establish an encrypted connection to Siam.AI's internal infrastructure from external networks. In this project, FortiClient VPN is essential for accessing the development server, ensuring that only authenticated users can connect to computing resources and datasets.

2.2.10 WangchanBERTa-finetuned-sentiment

The sentiment-thai-text-model is a pre-trained Transformer model for Thai sentiment analysis, published on Hugging Face by SandboxBhh. It is a fine-tuned version of WangchanBERTa, a Thai language model developed by NECTEC, and is further adapted to improve performance on Thai social-media style texts. The model is trained on the Wisersight Sentiment Corpus which contains Thai social-media messages annotated as positive, neutral, negative (and question in the original corpus). In practice, the model takes Thai text as input and categorizes it into three sentiment classes: positive, neutral or negative. In this project, this model is used to automatically assign sentiment labels to Thai documents, providing structured targets for downstream analysis and quality monitoring within the data pipeline.

2.3 Existing Solutions and Products in the Market

In recent years, many commercial and open-source platforms have been developed to address data ingestion, transformation and analytics at scale. This section compares the proposed system with major existing solutions in terms of architecture, flexibility, cost and suitability for Thai AI training workloads.

2.3.1 Cloud Provider Data Platforms (AWS, GCP, Azure)

Amazon Web Services (AWS). AWS offers a broad ecosystem for building data pipelines, including Amazon S3 for data lake storage, AWS Glue for serverless ETL and AWS Lake Formation for data lake creation and governance. Glue focuses on managed ETL and data preparation tightly integrated with other AWS analytics services such as Athena and Redshift. The pros include fully managed services, rich ecosystem and strong governance features. Limitations for our use case include vendor lock-in, ongoing cloud costs that conflict

with Siam.AI’s goal of using its own GPU infrastructure, less tailored Thai-specific preprocessing and data residency considerations.

Google Cloud Platform (GCP). On GCP, data ingestion and processing can be built using Cloud Storage as a data lake, BigQuery as the warehouse and Dataflow for unified batch and streaming processing. GCP provides a highly scalable, managed stack with strong streaming support, but assumes data and compute live in the cloud and is less suited to Siam.AI’s requirement for on-prem, open-source, Thai-centric processing.

Microsoft Azure (Azure Data Factory). ADF is a cloud-based data integration and ETL service that allows visual construction of data pipelines, with many pre-built connectors and a serverless execution model. Like AWS/GCP offerings, it is optimized for general enterprise analytics, not specifically for Thai LLM corpus preparation.

2.3.2 Open-Source Data Ingestion and Integration Tools (NiFi, Airbyte)

Apache NiFi is an open-source platform designed to automate the movement of data between systems. It supports real-time ingestion, transformation and routing with a drag-and-drop web UI and strong data provenance tracking. However, its UI-driven model can be heavier to maintain for highly customized text-centric preprocessing where Python code and Airflow DAGs may be easier to version-control, review and reuse.

Airbyte and other ELT tools focus on providing ready-made connectors to dozens or hundreds of data sources, primarily for structured or semi-structured data replication into cloud warehouses. Airbyte would help for standardized database/SaaS ingestion but offers limited support for the unstructured NECTEC corpora, complex Thai text cleaning or PDF and OCR workflows that are central in this project.

Table 2.1 Comparison of existing data platforms vs. the proposed Siam.AI pipeline

Topic	AWS Glue / Lake Formation	GCP Dataflow + BigQuery	Azure Data Factory + Synapse	Proposed Siam.AI Pipeline
Deployment	Fully managed cloud	Fully managed cloud	Fully managed cloud	On-prem, self-hosted
Main focus	General ETL / data lake	Batch + streaming pipelines	Visual data integration	End-to-end Thai AI data prep
Thai-specific preprocessing	Not Thai-centric	Not Thai-centric	Not Thai-centric	Thai-character ratio, PyThaiNLP, Thai OCR/ASR
Cost	Pay-as-you-go cloud	Pay-as-you-go cloud	Pay-as-you-go cloud	Existing Siam.AI hardware
Vendor lock-in	High	High	High	Low

2.3.3 Positioning of the Proposed System

- **On-prem, open-source, GPU-ready:** Built on MinIO, PostgreSQL, Apache Airflow and Python libraries, fully deployable on Siam.AI’s existing H100 server cluster, avoiding public cloud lock-in and recurring costs.
- **Thai-centric AI data preparation:** Pipeline logic is explicitly designed around Thai corpora (NECTEC datasets, Thai character ratio filtering, Thai-optimized OCR, ASR and PyThaiNLP integration).
- **Focus on AI-ready formats and metadata:** Outputs are standardized into JSONL, Parquet, TFRecords and backed by a rich metadata schema directly aligned with LLM training and evaluation needs.
- **Research and extensibility orientation:** Unlike fully managed cloud products, the system is intentionally transparent and code-driven, suitable for experimentation, benchmarking and future extension.

The proposed system therefore positions itself as a specialized, open-source, Thai-centric data ingestion and preparation pipeline, combining ideas from existing solutions but adapting them to the constraints and goals of Siam.AI’s internal AI model training platform.

2.4 Preliminary Design Based on Backgrounds

The preliminary design of the proposed system is based on the need to create an automated data ingestion and preparation pipeline for AI model training. The system is designed to support different types of input data, including audio files, PDF documents, JSON files and CSV files. These data types were selected because they represent common formats found in real-world AI dataset preparation tasks. However, raw data from these sources usually cannot be used directly for AI training because it may contain inconsistent formats, duplicated content, missing values, unreadable characters or unstructured information. Therefore, the system must provide a complete workflow for uploading, extracting, cleaning, validating, storing and monitoring data.

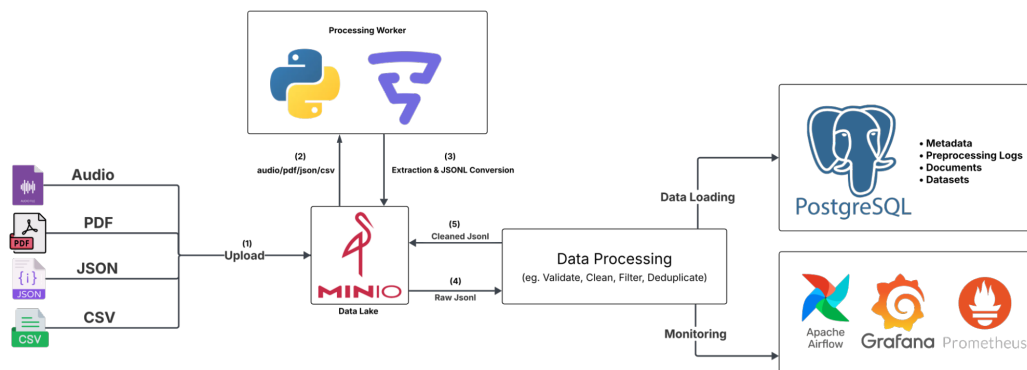


Figure 2.1 Automated Data Ingestion and Pipeline System Architecture

As shown in Figure 2.1, the architecture starts from the input layer, where users or data providers upload raw files into the system. The supported input formats include audio, PDF, JSON and CSV. These files are uploaded into MinIO, which acts as the main Data Lake of the system. MinIO is responsible for storing the original raw files as well as the cleaned outputs produced after processing. Using MinIO allows the system to organize data into different storage zones and maintain the separation between raw data and processed data.




Figure 2.2 NVIDIA B200 GPUs

After the raw files are uploaded into MinIO, the Processing Worker retrieves the input files for extraction and conversion. The Processing Worker is implemented using Python-based scripts and is orchestrated by Apache Airflow. Apache Airflow is used to manage the execution order of the pipeline tasks, schedule workflows, monitor task status and retry failed tasks when necessary.




Figure 2.3 MinIO Object Storage (Data Lake Interface)

The Processing Worker performs different tasks depending on the input file type. For PDF files, the system extracts text from documents. If the PDF contains embedded text, the text can be extracted directly. If the PDF is image-based or scanned, the system applies OCR to convert the document image into machine-readable text. For audio files, the system applies an ASR model to convert speech into text. In this project, audio processing is supported by the company's GPU infrastructure using NVIDIA B200, which provides high computational performance for model inference tasks. For JSON and CSV files, the system reads the structured or semi-structured content and converts it into a common JSON-based format for further processing.

After extraction and conversion, the output from the Processing Worker is converted into raw JSON. This raw JSON acts as an intermediate representation of the extracted data and helps standardize different sources before they enter the data processing layer. The raw JSON is then passed to the Data Processing layer, which is responsible for improving the quality of the extracted data before it is stored as final output. The main operations in this layer include validation, cleaning, filtering and deduplication.

This processing layer is especially important for Thai-language datasets, which may include encoding issues, inconsistent spacing, missing tone marks or OCR/ASR errors. For PDF data, this may include removing low-quality OCR results or empty pages. For audio data, this may include filtering inaccurate transcripts or silent audio. For JSON and CSV data, this may include checking required fields, normalizing text and removing duplicate records. Once the data has passed the processing stage, the cleaned result is stored back into MinIO as cleaned JSONL.




Figure 2.4 pgAdmin 4 Interface for PostgreSQL Data Warehouse

In addition to storing cleaned outputs in MinIO, the system also loads important information into PostgreSQL. PostgreSQL acts as the main database for structured storage, querying and metadata management. The database stores key information such as datasets, documents, preprocessing logs and metadata. The connection between the data processing layer and PostgreSQL is important because AI-ready datasets should not only contain cleaned text but also processing history and metadata.




Figure 2.5 Apache Airflow Interface for Data Pipeline Connection

The final part of the architecture is the Monitoring layer. The system uses Apache Airflow, Grafana and Prometheus to monitor pipeline execution and system performance. Airflow provides task-level monitoring, while Prometheus collects metrics and Grafana visualizes them through dashboards. These monitoring tools help the development team track whether the pipeline is running correctly and identify problems such as failed tasks, slow processing, high resource usage or abnormal data volume.

Overall, the updated architecture follows a clear end-to-end flow: raw input files are uploaded into MinIO; the Processing Worker extracts and converts the files into raw JSON; the Data Processing layer validates, cleans, filters and removes duplicates; the cleaned JSONL output is stored back into MinIO; metadata is loaded into PostgreSQL; and Airflow, Prometheus and Grafana are used to monitor pipeline execution and system health. This design focuses on practical batch-based data ingestion and preparation rather than real-time streaming.

CHAPTER 3 METHODOLOGY AND DESIGNS

3.1 Introduction

This chapter presents the current progress of methodology and system design for the development of the Automated Data Ingestion and Pipeline System for AI Model Training. The methodology is grounded in industry-proven data engineering principles and aligns with the project scope, requirements and metadata schema. This chapter details the functional requirements, the overall architecture, the containerized service design, data lake and data warehouse structures, the entire pipeline orchestration mechanisms, and the monitoring framework used to ensure reliability and scalability.

3.2 Functional Requirements

The functional requirements were derived from the system objectives, the dataset schema and the expected output formats for AI model training. These requirements define what the system must accomplish to support fully automated data ingestion and preparation.

3.2.1 Data Ingestion Requirements

The system is capable of ingesting the following data types:

- Structured data: CSV, SQL, JSON, JSONL
- Unstructured data: raw text, scanned PDFs, images and audio

3.2.2 Data Cleaning and Transformation Requirements

- Perform encoding normalization, whitespace and punctuation cleaning and structural standardization.
- Remove duplicate data based on checksums as defined in the metadata schema.
- Apply Thai-specific preprocessing rules including Thai-character ratio filtering required for NECTEC corpora.
- Normalize and flatten formats into standardized JSONL or Parquet representations.
- Extract text from PDF and image data using an OCR engine optimized for Thai language.
- Extract text from audio files using an ASR engine optimized for Thai language.

3.2.3 Output and Storage Requirements

- Cleaned intermediate data in a clean zone.
- Final AI-ready datasets in processed zones using formats such as JSONL, Parquet and TFRecords.
- Metadata stored in a relational warehouse, following the provided schema definitions.

3.2.4 Orchestration and Monitoring Requirements

- Automate all ingestion, cleaning, OCR and ASR workflows through Apache Airflow DAGs.
- Provide traceability and fault tolerance through logging, retries and status tracking.
- Incorporate Prometheus and Grafana for system-level monitoring as required in the architecture plan.

3.3 System Design

3.3.1 Architectural Overview

The system follows a modular Data Lakehouse architecture composed of:

- A Data Lake (MinIO) for storing raw, cleaned and processed files.
- A Data Warehouse (PostgreSQL) for metadata management and querying.
- An Orchestration Layer (Apache Airflow) for pipeline automation.
- Python-based preprocessing modules for structured, semi-structured and unstructured data.
- Monitoring services (Prometheus, Grafana) for system observability.

3.3.2 Containerized Service Architecture

All system components are deployed using the company's platform to ensure reproducibility, isolation and consistent execution across environments. Each service plays a specific role within the overall pipeline:

- **MinIO** (minio/minio:latest) serves as the Data Lake for all raw, cleaned and processed data. Its S3-compatible API supports seamless integration with Python scripts, Airflow tasks and machine learning frameworks.
- **PostgreSQL** (postgres:13) implements the Data Warehouse using the metadata schema.
- **Airflow** Webserver, Scheduler and Init (apache/airflow:3.1.0) implement DAGs.
- **pgAdmin** (dpage/pgadmin4) provides administrative access to the metadata warehouse.
- **Prometheus** (prom/prometheus) collects metrics from exporters (Airflow, PostgreSQL, MinIO) and stores them as time-series data.
- **Grafana** (grafana/grafana) reads metrics from Prometheus and provides dashboards.
- **PostgreSQL Exporter** (prometheuscommunity/postgres-exporter:latest) exposes database metrics to Prometheus.
- **StatsD Exporter** (prom/statsd-exporter:v0.26.0) receives StatsD metrics from Airflow.

3.4 Data Lake Structure




Figure 3.1 MinIO Interface

3.4.1 MinIO Structure

MinIO is used as the object storage layer for all datasets in the pipeline. A main bucket is created and organized into logical prefixes that follow the processing steps of the pipeline: `raw`, `cleaned`, `validated`, `meta` and `mapping`. The `raw` prefix stores files exactly as they are received from the original sources. The `cleaned` prefix contains the output of the text cleaning scripts, while the `validated` prefix stores only the files that pass schema and data-quality checks and are ready to be used for model training or downstream analytics. The `meta` prefix is used to store log files from each processing step. The `mapping` prefix stores the outputs from an additional mapping step required for some NECTEC datasets, in which raw files are joined with CSV mapping files provided by NECTEC and saved for later ingestion into the database. This folder structure makes it clear which stage each file belongs to and supports a simple, traceable flow from raw data to validated data.

3.4.2 Bucket Versioning and Replication

Bucket versioning is enabled for the main bucket so that previous versions of objects are kept when files are updated or overwritten during pipeline re-runs. This helps protect against accidental data loss and allows the system to roll back to an earlier version of a file if needed. In addition, lifecycle rules are configured to automatically clean up old object versions and temporary files after a retention period, which prevents unlimited growth of storage usage. MinIO bucket replication is also configured so that objects written to the main bucket are asynchronously copied to a second bucket on another MinIO deployment. This replication provides an additional backup of the cleaned and validated data and improves the resilience of the overall system.

3.5 Data Warehouse




Figure 3.2 Database schema for PostgreSQL

The data warehouse uses PostgreSQL as a central database for managing dataset metadata, documents, labels and processing logs. The schema is designed around a small set of core tables with UUID primary keys and foreign-key links between them.

Table 3.1 datasets table schema

Field name	Type	Description
id	UUID / BIGINT	Primary key
name	VARCHAR(255)	Dataset name
description	TEXT	Dataset details
source_url	TEXT	Source URL or repository
domain	VARCHAR(50)	Categories such as general, medical, legal, tourism
language	VARCHAR(10)	Main languages of the dataset
version	VARCHAR(20)	Dataset version
created_at	TIMESTAMP	Creation date
updated_at	TIMESTAMP	Last edit date
status	VARCHAR(20)	raw, preprocessed, validated

Table 3.2 documents table schema

Field name	Type	Description
id	UUID / BIGINT	Primary key
dataset_id	UUID	Foreign key → datasets.id
parent_id	UUID	Reference to main document if sub-document
title	VARCHAR(255)	Topic / filename
content	TEXT	All content
content_tokens	JSONB	Tokenized version (optional)
language	VARCHAR(10)	th, en, cn, etc.
type	VARCHAR(50)	paragraph, code, conversation, qa
source	VARCHAR(255)	Source
checksum	VARCHAR(64)	Hash of content to prevent duplicates
word_count	INT	Word count
char_count	INT	Character count
metadata	JSONB	Topic, domain, difficulty
created_at	TIMESTAMP	Creation date
updated_at	TIMESTAMP	Last edit date
status	VARCHAR(20)	raw, cleaned, validated

Table 3.3 conversations table schema

Field name	Type	Description
id	UUID	Primary key
dataset_id	UUID	Foreign key → datasets.id
conversation_id	VARCHAR(100)	Identifier of conversation group
title	VARCHAR(255)	Name or topic of conversation
domain	VARCHAR(50)	Category, e.g., customer_service, qa
metadata	JSONB	e.g., intent, scenario
created_at	TIMESTAMP	Creation date
updated_at	TIMESTAMP	Last edit date

Table 3.4 messages table schema

Field name	Type	Description
id	UUID	Primary key
conversation_id	UUID	Foreign key → conversations.id
parent_message	UUID	Reference to main topic if reply
role	VARCHAR(20)	user, assistant, system, moderator
content	TEXT	Information
content_tokens	JSONB	Tokenized version
length_tokens	INT	Token count
sentiment	VARCHAR(20)	e.g., positive, negative, neutral
entities	JSONB	e.g., {"name": xyz}
intent	VARCHAR(100)	Optional for fine-tuning
metadata	JSONB	Additional info, e.g., topic, urgency
timestamp	TIMESTAMP	Created message
status	VARCHAR(20)	raw, cleaned, validated

Table 3.5 labels table schema

Field name	Type	Description
id	UUID	Primary key
target_id	UUID	Foreign key → messages.id or documents.id
label_type	VARCHAR(50)	e.g., intent, topic, sentiment, category
label_value	VARCHAR(255)	Label value, e.g., greeting, legal, positive
confidence	DECIMAL(5,2)	Confidence score (0–100%)
annotator	VARCHAR(100)	User or AI model that created the label
created_at	TIMESTAMP	Creation date

Table 3.6 preprocessing_logs table schema

Field name	Type	Description
id	UUID	Primary key
target_id	UUID	Foreign key → documents.id or messages.id
step	VARCHAR(50)	e.g., remove_html, normalize_unicode, tokenize
status	VARCHAR(20)	success, failed
details	TEXT	Result info, e.g., error message
executed_by	VARCHAR(50)	Script or user
executed_at	TIMESTAMP	Processed time

Table 3.7 evaluation_metrics table schema

Field name	Type	Description
id	UUID	Primary key
dataset_id	UUID	Foreign key → datasets.id
model_name	VARCHAR(100)	Model name
metric_type	VARCHAR(50)	e.g., accuracy, bleu, rouge, perplexity
metric_value	DECIMAL(10,4)	Metric value
evaluated_at	TIMESTAMP	Evaluation date
notes	TEXT	More info, e.g., subset, conditions

Relationships:

1. datasets $\rightarrow$ documents (1:N)
2. datasets $\rightarrow$ conversations (1:N)
3. conversations $\rightarrow$ messages (1:N)
4. documents/messages $\rightarrow$ labels (1:N)
5. documents/messages $\rightarrow$ preprocessing_logs (1:N)

The `datasets` table stores one record for each dataset used in the pipeline. The `documents` table stores individual text units that are ready to be used for analysis or model training. Each document row is linked to a dataset through `dataset_id`. The `labels` table is a generic table for storing labels for any target record, allowing the system to attach many different kinds of labels to the same text without changing the main tables. The `preprocessing_log` table stores detailed logs for each processing step applied to a given target.

3.6 Orchestration

3.6.1 JSONL Pipeline

The end-to-end JSONL pipeline is orchestrated using Apache Airflow in a DAG named `sensor_demo`. This DAG is scheduled to run every hour and is designed to automatically detect new raw `.jsonl` files in MinIO, run the full preprocessing pipeline and update both the PostgreSQL database and the monitoring logs. Whenever new files appear in the configured MinIO prefix, Airflow picks them up and processes them without manual intervention.




Figure 3.3 Apache Airflow graph view DAGs

The first section of the DAG focuses on detecting new and stable input files. The task `wait_raw_stable` uses `S3KeysUnchangedSensor` to make sure that files in the `raw` prefix are no longer changing before starting the pipeline. After that, `list_raw_keys` uses `S3ListOperator` to list all raw object keys under the configured bucket and prefix. The next task, `choose_keys`, is a Python `@task` that filters this list according to the dataset and directory structure. Then, `filter_new` compares the selected keys with previously processed files and keeps only those that have not yet been ingested.

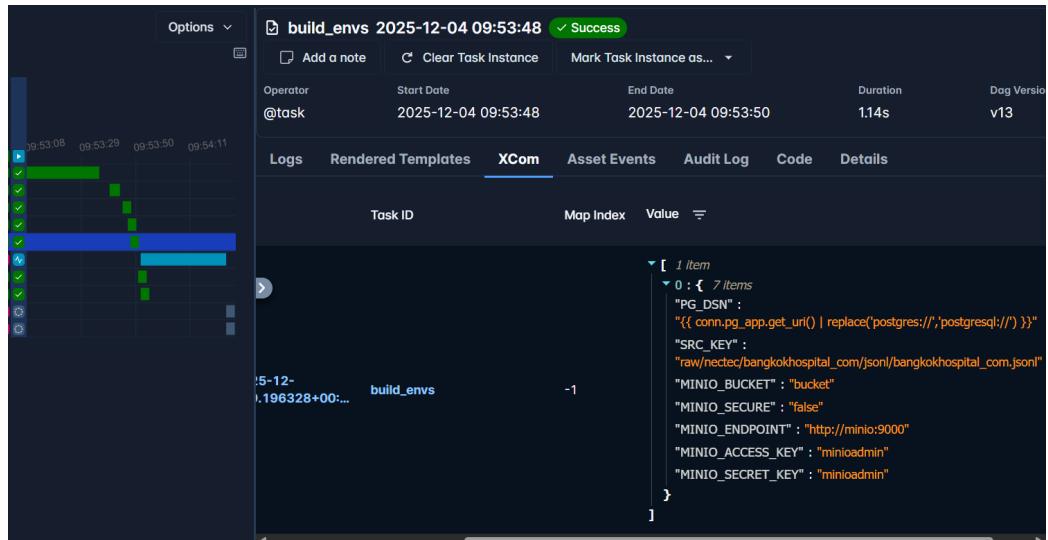


Figure 3.4 Apache Airflow Build environment task detail

Once the DAG has a list of new raw files, the `build_envs` task prepares the runtime configuration for each file. This task constructs environment variables such as source key, dataset name, and various MinIO and database parameters required by the downstream scripts. The DAG then uses Airflow's dynamic task mapping to create one set of downstream tasks per input file.

After `build_envs`, the pipeline splits into three parallel branches that handle ingestion, logging and labelling. The `run_pipeline_all` task is a `BashOperator` that calls the main `run.py` script for each file, which performs the core operations: reading the raw `.jsonl` from MinIO, cleaning and validating the text, running any required mapping step, writing the cleaned or mapped outputs back to MinIO, and inserting document records into the PostgreSQL documents table. In parallel with ingestion, the `to_log_params` task prepares log parameters for each processed file, and the `log-ingest` task writes a structured log entry into the `preprocessing_log` table.

The third branch handles automatic labelling. The `to_label_params` task prepares the parameters required to call the Thai sentiment model. The `sentiment_upsert_mapped` task then runs a Python function that loads the documents, calls the sentiment model and upserts the results into the `labels` table.




Figure 3.5 Apache Airflow Sensor_demo DAG detail

At the current stage, the `sensor_demo` DAG is stable and has been executed many times, with most runs marked as successful. The DAG runs on a schedule and can also be triggered manually for testing. New raw files uploaded to the MinIO `raw/` prefix are correctly detected by the sensor, passed through the dynamic mapping pipeline and result in new documents, labels and `preprocessing_log` records in PostgreSQL.

3.6.2 Unstructured PDF OCR Pipeline

The PDF OCR pipeline is orchestrated using Apache Airflow in a DAG named `pdf_ocr_pipeline`. This DAG is scheduled to run daily and is designed to automatically detect new PDF files in MinIO, run a hybrid text-extraction pipeline against them, and write both the clean text outputs and the per-page processing logs back to MinIO. Whenever new PDFs appear in the configured `raw/` prefix, or whenever a previous run produced documents marked as low-quality or failed, Airflow will pick them up and re-process them without manual intervention.




Figure 3.6 Apache Airflow graph view of `pdf_ocr_pipeline`

The first section of the DAG focuses on discovering and filtering PDF files. The `find_pdfs` task scans every bucket under the configured `raw/` prefix and lists all `.pdf` objects, while skipping any key that already sits under a `jsonl/` subdirectory. For each candidate file, the task reads the existing result and meta JSONL files in MinIO to check whether the PDF has already been processed successfully, or whether it previously ended in a `low_quality`, `empty` or partially failed state. Documents that have already been processed cleanly are skipped, while those that failed or were flagged as low quality are re-queued for another attempt up to a configurable cap (`MAX_RUN_ATTEMPTS = 2`). This makes the pipeline idempotent and lets it recover gracefully from transient OCR-server failures across runs.

Once the list of PDFs is built, the DAG uses Airflow's dynamic task mapping to create one `process_pdf` task instance per file. Each mapped task runs inside a cached virtual environment (`/tmp/pdf_ocr_venv`) that ships with all required libraries such as PyMuPDF, Pillow, NumPy, PyThaiNLP, requests and boto3, and is bound to a dedicated Airflow pool named `ocr_pool` with a per-DAG-run concurrency limit of four. This

configuration prevents the pipeline from overwhelming the upstream Typhoon-OCR service running on the B200 GPU.

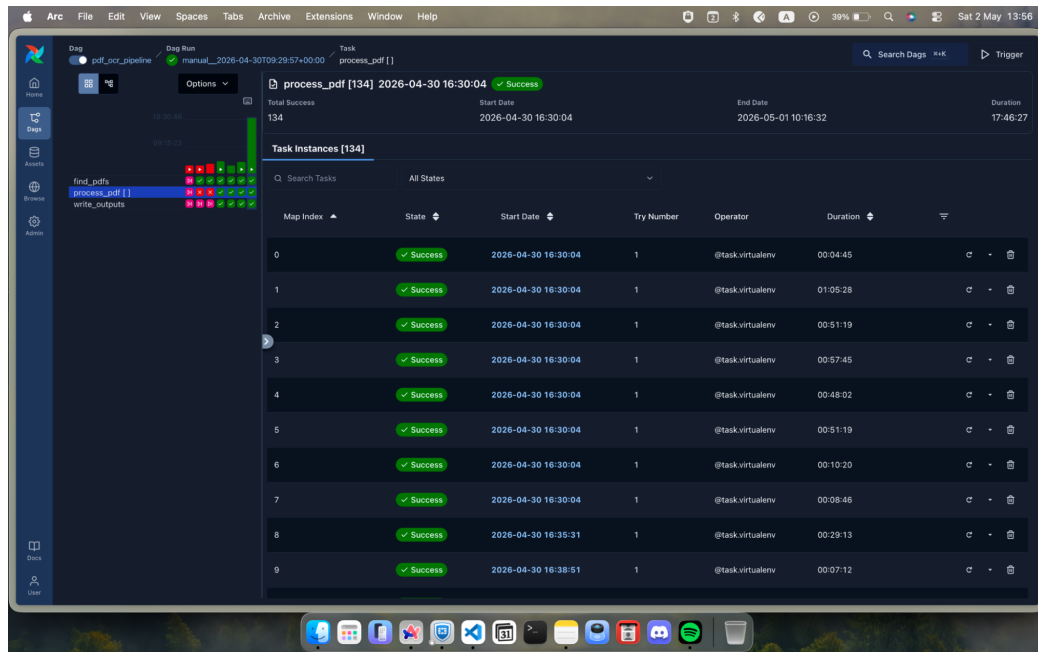


Figure 3.7 Apache Airflow process_pdf task detail

Inside each `process_pdf` task, the file is downloaded from MinIO and opened with PyMuPDF. The pipeline first attempts native text extraction at the page level: if a page yields enough characters of valid Thai text and shows no signs of font-encoding failure (such as `(cid:N)` glyphs, replacement characters or unmapped private-use-area codepoints), its text is kept and the page is recorded as having been handled by the PyMuPDF method. Pages that fail this check are routed to the OCR fallback path. Before invoking the OCR server, the task performs a health check; if the server is reachable, each remaining page is rendered to an RGB image at the configured DPI cascade (400 DPI primary, 300 DPI fallback when the JPEG exceeds the size limit), optionally preprocessed with grayscale conversion, contrast enhancement and deskewing, and then sent in parallel to the Typhoon-OCR endpoint together with anchor text taken from the PyMuPDF extraction. Returned text is filtered through a series of quality gates that detect common hallucination patterns, then cleaned to remove markdown formatting, commentary blocks and English image-description paragraphs that the model occasionally emits.

The final stage of the DAG is the `write_outputs` task, which is configured with `trigger_rule="all_done"` so that it executes even if some mapped tasks have failed. This task groups all returned records by bucket and folder and upserts them into two JSONL files per source folder: a result file containing only the successfully extracted documents, and a meta file containing detailed page-level logs.




Figure 3.8 Apache Airflow pdf_ocr_pipeline DAG run detail

At the current stage, the pdf_ocr_pipeline DAG is stable and runs reliably on its daily schedule. New PDF files uploaded to the MinIO raw/ prefix are correctly discovered, processed through the hybrid PyMuPDF + Typhoon-OCR pipeline, and produce both clean training-text outputs in the result JSONL and rich diagnostic information in the meta JSONL.

3.6.3 Unstructured Audio ASR Pipeline

The audio ASR pipeline is orchestrated using Apache Airflow in a DAG named `audio_asr_pipeline`. This DAG runs on a daily schedule and is designed to detect new audio files in MinIO, transcribe them with a Thai/Isan Whisper-based ASR service, and write both the clean transcripts and detailed quality metadata back to MinIO. The DAG supports two input formats: standalone audio files (`.wav`, `.mp3`, `.flac`, `.m4a`, `.ogg`, `.opus` and other common containers) and Parquet datasets in which each row carries an embedded audio payload together with an optional ground-truth transcription.




Figure 3.9 Apache Airflow graph view of `audio_asr_pipeline`

The first task, `find_audio_files`, lists every object under the configured `raw/` prefix across all buckets and selects files whose extensions match either the supported audio formats or `.parquet`. Once the list is built, the DAG uses dynamic task mapping to fan out into one `process_audio_file` instance per discovered file. Mapped tasks are bound to the dedicated `asr_pool` Airflow pool with up to eight concurrent instances per DAG run, and each task in turn spawns its own internal pool of eight worker threads, giving an effective fan-out of roughly thirty-two concurrent ASR requests against the GPU server.




Figure 3.10 Apache Airflow `process_audio_file` task detail

Inside each `process_audio_file` task, the pipeline first checks the health of the ASR endpoint and reads any previously generated result and meta JSONL files so that already-transcribed rows can be skipped. The processing logic then branches on the file type. For standalone audio files, the task downloads the bytes from MinIO, measures the duration and either transcribes the file in a single request or, if the audio exceeds the configured split threshold of 300 seconds, decodes it to 16 kHz mono with PyAV and splits it into approximately 30-second WAV segments that are transcribed independently and then merged into a single record. For Parquet inputs, the task iterates over the file in batches using `pyarrow`, extracts the audio bytes and any ground-truth column from each row, and dispatches the rows to a `ThreadPoolExecutor`.

Each transcription is then evaluated by a quality function. When ground truth is available (Path A), the function computes the Character Error Rate between the hypothesis and the reference after applying Thai-specific normalization. If the initial CER exceeds the configured threshold, the function attempts a series of recovery strategies: stripping trailing Whisper repetition loops, expanding known abbreviations on the ground-truth side, retrying decoding at a higher temperature and applying a foreign-word recovery rule. When ground truth is missing (Path B), the function instead checks that the output is non-trivial, free of repetition loops and consistent with a secondary higher-temperature transcription on sufficiently long clips. Successful rows are written to the result JSONL, while low-quality, silent or failed rows are written to the meta JSONL together with a quality reason, the measured duration, the CER and abbreviation hints when applicable.




Figure 3.11 Apache Airflow `audio_asr_pipeline` DAG run detail

After all chunks have been processed, the task performs a final deduplication pass over both output JSONL files using last-write-wins semantics keyed on filename, records per-file quality statistics and returns this summary to the DAG. The downstream `write_outputs` task aggregates these per-file summaries into a global report.

3.7 Monitoring

For this project, a monitoring system was implemented using Prometheus and Grafana to track the health and performance of the whole data pipeline. All monitoring services run as Docker containers in the same environment as Airflow, MinIO and PostgreSQL.

3.7.1 Metrics Collection with Prometheus

Prometheus is used as the metrics collector. It is configured through a `prometheus.yml` file, where each component is defined as a separate scrape job with its endpoint and port. Prometheus regularly calls the metrics endpoint of each exporter and stores the results in its time-series database.




Figure 3.12 Prometheus metric connector

- **Airflow exporter and airflow_statsd** export DAG and task metrics from the Airflow webserver and scheduler. These metrics include DAG run status, task duration, number of retries and scheduler activity.
- **MinIO and minio-bucket** exporters expose metrics from the object storage. One focuses on cluster status (disk usage, read/write operations, latency), while the bucket exporter reports bucket-level information.
- **Postgres exporter** reads internal statistics from PostgreSQL and exposes metrics like query duration, number of connections, table size and cache hit rate.
- **Prometheus** itself is also scraped, allowing basic self-monitoring.

In the Prometheus “Targets” page, all these jobs appear as separate targets, and the status “1/1 up” confirms that Prometheus can successfully connect to each exporter and collect metrics on a regular schedule.

3.7.2 Grafana Dashboard

Grafana is connected to Prometheus as a single data source by pointing it to the Prometheus HTTP endpoint. All dashboards use PromQL queries to read metrics from Prometheus and display them as graphs, tables and single-value panels. To keep things organized, the dashboards are grouped into folders by component.




Figure 3.13 Grafana Dashboard

Overall, this monitoring setup gives a single place to see whether the pipeline is running, how long it takes, how much data is stored and whether the data-quality checks behave as expected. This information is later used as the basis for alert rules so that problems in Airflow, MinIO or PostgreSQL can be detected early.

CHAPTER 4 RESULTS AND DISCUSSION

4.1 Data Lake

The MinIO bucket `ai-datasets` is organized into multiple layers to support different stages of the data processing pipeline. Within the `raw/` directory, datasets are further organized by source. This structure ensures that data from multiple domains and sources can be managed efficiently and allows easy extension when new datasets are added to the system. The organization aligns with the dataset catalog used in the project, which includes multiple domains such as finance, medical, tourism and legal.




Figure 4.1 Datalake Layer




Figure 4.2 NECTEC Datasets in MinIO

4.2 Unstructured Pipeline

The unstructured pipeline was developed to process data that does not already exist in a structured table format. In this project, the unstructured pipeline consists of two main components: the PDF pipeline and the audio pipeline. The PDF pipeline extracts Thai text from both text-native and image-heavy PDF files, while the audio pipeline converts Thai speech into text using an Automatic Speech Recognition (ASR) model.

Table 4.1 Overall unstructured pipeline output

Pipeline	Records / Documents	Output volume
PDF pipeline	133 documents / 7,346 pages	8.42 million characters
Audio pipeline	307,793 utterances	19.42 million characters / 656.93 audio hours
Total	—	27.85 million characters

In conclusion, the unstructured pipeline successfully converted both PDF and audio data into AI-ready text outputs. The PDF pipeline extracted 8.42 million characters from Thai PDF documents, while the audio pipeline generated 19.42 million characters from Thai speech data. Together, both pipelines produced approximately 27.85 million characters of usable text data.

4.2.1 PDF Pipeline

The PDF pipeline was evaluated using a Thai-language PDF corpus from NECTEC. The dataset was divided into two sub-corpora: pdf1 and pdf2. The pdf1 dataset mainly contains text-native PDFs, while pdf2 mainly contains image-heavy or scanned PDF documents. In total, the pipeline processed 133 PDF documents containing 7,346 pages and 1,211.6 MB of data.

Table 4.2 PDF dataset overview

Dataset	Document type	Documents	Pages	Size (MB)
pdf1	Text-native PDFs	94	5,181	693.6
pdf2	Image-heavy PDFs	39	2,165	518.1
Total	Mixed PDF corpus	133	7,346	1,211.6

At the document level, the pipeline successfully processed all PDF files. Both pdf1 and pdf2 achieved a 100% document success rate, with no failed documents.

Table 4.3 PDF document-level results

Metric	pdf1	pdf2	Combined
Documents processed	94	39	133
Documents succeeded	94 (100%)	39 (100%)	133 (100%)
Documents failed	0	0	0
Total pages	5,181	2,165	7,346
Average pages / document	55.1	55.5	55.2
Total characters extracted	5,823,040	2,601,378	8,424,418
Total processing time	174.5 min	152.9 min	327.4 min
Average time / document	111.4 sec	235.2 sec	—

Although pdf2 contained fewer documents than pdf1, it required more time per document because most pages were image-heavy and needed OCR processing. This shows that scanned or image-based PDFs require more computational resources than text-native PDFs.

Table 4.4 PDF page-level results

Page status	pdf1	pdf2	Combined
Success	4,960 (95.7%)	2,032 (93.9%)	6,992 (95.2%)
Low quality	11 (0.2%)	11 (0.5%)	22 (0.3%)
Empty	210 (4.1%)	122 (5.6%)	332 (4.5%)
Failed	0	0	0

The PDF pipeline used a hybrid extraction method. It used PyMuPDF as the primary text extractor and applied OCR only when direct text extraction was not sufficient. This approach improved efficiency because text-native pages could be processed quickly without OCR, while image-heavy pages could still be recovered through OCR.

Table 4.5 PDF extraction method results

Metric	pdf1	pdf2
Pages routed to PyMuPDF	3,730 (72.0%)	580 (26.8%)
Pages routed to OCR	1,451 (28.0%)	1,585 (73.2%)
OCR page-success rate	1,230 / 1,451 (84.8%)	1,452 / 1,585 (91.6%)
Character share from PyMuPDF	5,312,126 (92.9%)	617,743 (24.6%)
Character share from OCR	408,645 (7.1%)	1,892,446 (75.4%)
Documents needing OCR only	0	22 (56.4%)
Documents using both methods	94 (100%)	17 (43.6%)

The results show that the hybrid method was important for both datasets. In pdf1, most text was extracted using PyMuPDF because the documents were mainly text-native. In pdf2, most text came from OCR because the documents were image-heavy. The PDF pipeline achieved reliable results for both text-native and image-heavy PDF files. Key findings:

- The pipeline achieved 100% document-level success.
- The combined page success rate was 95.2% with no failed pages.
- The hybrid extraction strategy was effective.
- The pdf2 dataset required more OCR processing because it contained more scanned and image-heavy layouts.
- The low-quality rate was only 0.3%, showing that the quality gate successfully removed unreliable outputs.
- Empty pages were excluded intentionally to reduce the risk of hallucinated text entering the final training dataset.

4.2.2 Audio Pipeline

The audio pipeline was evaluated using a NECTEC Thai-central audio corpus. The dataset was stored as 16 Parquet shards, each containing audio bytes with corresponding ground-truth Thai transcripts. The pipeline used the Typhoon Isan ASR Whisper model to convert speech into text. In total, the pipeline processed 335,674 audio rows.

Table 4.6 Audio dataset overview

Metric	Value
Dataset source	NECTEC Thai-central audio corpus
Number of shards	16 Parquet shards
Total rows processed	335,674
Model used	typhoon-ai/typhoon-isan-asr-whisper
Main quality metric	Character Error Rate (CER)
Acceptance threshold	$\text{CER} \leq 0.20$

Table 4.7 Audio acceptance results

Outcome	Count	Percentage
Kept / Success	307,793	91.69%
Filtered as low quality	27,260	8.12%
Filtered as silent audio	621	0.19%
Failed / server empty	0	0.00%
Total	335,674	100.00%

The pipeline kept 307,793 records, representing 91.69% of all input rows. This high acceptance rate shows that most audio samples were successfully transcribed with acceptable quality. The pipeline also produced zero failed or server-empty records, indicating that the ASR workflow was stable.

Table 4.8 Audio output volume

Metric	Value
Kept transcripts	307,793 records
Total audio kept	656.93 hours
Total audio kept	2,364,958 seconds
Total transcript characters	19,423,129
Average characters / utterance	63.1

The accepted records showed strong transcription quality. The mean CER was 0.0447 and the median CER was 0.0270. In addition, 38.2% of accepted utterances had a perfect CER of 0, meaning the predicted transcription exactly matched the ground-truth transcript.

Table 4.9 CER distribution of accepted records

Statistic	Value
Mean CER	0.0447
Median CER	0.0270
25th percentile	0.0000
75th percentile	0.0741
90th percentile	0.1250
Records with CER = 0	117,619 (38.2%)
Records with CER < 0.05	196,716 (63.9%)
Records with CER < 0.10	256,590 (83.4%)

Table 4.10 CER distribution of filtered records

Statistic	Value
Mean CER	0.3618
Median CER	0.2909
Minimum CER	0.2007
Maximum CER	11.0857
Quality reason	high_cer

The rejected records likely included incorrect transcriptions, noisy audio, wrong-language outputs or other ASR errors. Most accepted audio clips were short utterances. The median duration was 6.40 seconds, while the mean duration was 7.68 seconds.

Table 4.11 Audio duration profile of accepted records

Metric	Duration
Mean duration	7.68 sec
Median duration	6.40 sec
90th percentile	14.51 sec
99th percentile	21.16 sec
Maximum duration	61.18 sec

Table 4.12 Per-shard result consistency

Metric	Range across 16 shards
Kept rows	19,138 – 19,278
Average CER of accepted records	0.0441 – 0.0451
Audio hours kept	40.66 – 41.30
Filtered low-quality records	1,661 – 1,803
Silent audio records	28 – 48

The variation across shards was less than 1% for the main metrics. To verify that the audio pipeline was not limited to a specific audio container or codec, an audio format compatibility test was conducted. The same Thai news audio clip, named `sukie-test`, was encoded into 13 different audio formats and processed through the full ASR pipeline.

Table 4.13 Audio format compatibility test results

No.	Format	Container / Codec family	Status	Transcript chars
1	.wav	PCM, uncompressed	Success	259
2	.flac	Free Lossless Audio Codec	Success	259
3	.aiff	Apple uncompressed	Success	259
4	.au	Sun/NeXT μ -law	Success	259
5	.mp3	MPEG-1/2 Layer III	Success	259
6	.aac	Advanced Audio Coding	Success	258
7	.m4a	MP4 container / AAC audio	Success	259
8	.mp4	MP4 container, audio track	Success	259
9	.ogg	Ogg Vorbis	Success	259
10	.opus	Opus	Success	259
11	.webm	WebM container / Opus	Success	259
12	.wma	Windows Media Audio	Success	258
13	.3gp	3GPP, mobile	Success	259

The result shows that 13 out of 13 formats were successfully decoded and transcribed, giving a 100% format success rate. This is important because audio data collected from different sources may not always follow a single standard format. To verify that the result was not dependent on only one source clip, the same test was repeated using a different short 6.4-second Thai utterance: “ทีมจากอิสราเอลไม่ควรได้เป็นเจ้าบ้านในเกมส์ฟุตบอล”.

Table 4.14 Format compatibility replication results

Format group	Formats	Status	Chars	Transcript result
Lossless / uncompressed	.wav, .flac, .aiff, .au	Success	49	Identical Thai sentence
MP / AAC family	.mp3, .aac, .m4a, .mp4	Success	49	Identical Thai sentence
Open web formats	.ogg, .opus, .webm	Success	49	Identical Thai sentence
Legacy / mobile formats	.wma, .3gp	Success	49	Identical Thai sentence

Table 4.15 Combined format-coverage summary

Test set	Source clip	Duration (s)	Formats	Successful	Byte stability
test_audio_formats	Thai news clip	~30	13	13/13 (100%)	11/13 (84.6%)
test_audio	Thai single sentence	6.4	13	13/13 (100%)	13/13 (100%)
Combined	—	—	26 trials	26/26 (100%)	24/26 (92.3%)

The audio pipeline achieved strong results in transforming Thai speech data into AI-ready text transcripts.

Key findings:

- The pipeline accepted 91.69% of all audio records after applying the CER quality gate.
- The final accepted set contained 656.93 hours of audio and 19.42 million transcript characters.
- The mean CER of accepted records was only 0.0447.
- 38.2% of accepted records had a perfect CER of 0.
- The pipeline filtered out 8.12% low-quality records.
- Silent-audio detection removed 0.19% of records that did not contain meaningful speech.
- No failed or server-empty records appeared in the final output.
- The format compatibility tests showed 100% success across 26 format-clip trials.

One important reason for using CER instead of WER is that Thai language does not always separate words with spaces. CER is more suitable for this project because it evaluates transcription quality at the character level.

4.3 JSONL Pipeline

4.3.1 Main Pipeline Execution

The main pipeline demonstrates a successful execution of the end-to-end workflow. The DAG consists of multiple tasks that handle different stages of the data ingestion and preparation process. Key steps in the pipeline include:

- `wait_raw_stable`: waits for new data to be available in the storage system
- `list_raw_keys`: retrieves file paths from the data lake (MinIO)
- `choose_keys`: selects relevant datasets for processing
- `filter_new`: filters out already processed or duplicate data
- `build_envs`: prepares runtime environment and parameters
- `run_pipeline_all`: executes the main processing logic
- `to_label_params` / `to_log_params`: prepares parameters for labeling and logging
- `sentiment_upsert_mapped`: stores processed results into the database
- `log_ingest`: logs ingestion results

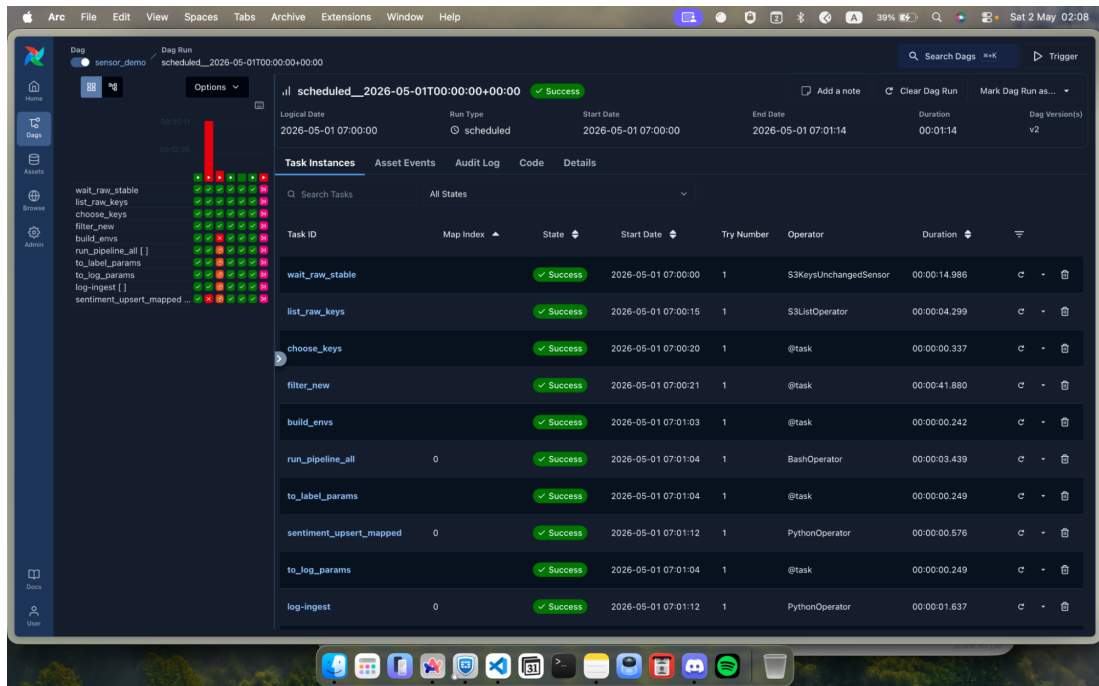


Figure 4.3 Main pipeline DAG

4.3.2 Sentiment Processing Pipeline

The second pipeline focuses on sentiment analysis and labeling of processed text data. The task `run_sentiment_from_db` executes successfully using a `PythonOperator`. This demonstrates that the pipeline supports incremental processing and can continuously enrich datasets with labels, which is essential for supervised learning and model fine-tuning.




Figure 4.4 Sentiment Processing Pipeline

4.4 Database

This section presents the implementation and results of the PostgreSQL database used in the data ingestion and preparation pipeline. The database is designed to support structured storage of datasets, documents, labels and preprocessing logs, enabling efficient querying and monitoring of data processing workflows. The schema is designed with a relational structure consisting of multiple interconnected tables including datasets, documents, labels and preprocessing_logs, which together support the full lifecycle of data ingestion, transformation and validation.




Figure 4.5 Database schema in PostgreSQL

4.4.1 Datasets

The datasets table contains metadata describing each dataset ingested into the system. This includes dataset name, domain, language and version. This demonstrates that the pipeline is capable of handling diverse data sources and domains, which is essential for building generalized AI models. The dataset categorization also enables filtering and selection of data for specific use cases.

The screenshot shows the PGAdmin interface with the 'public.datasets' table selected. The table contains 22 rows of data. The columns are: id (PK), name (character varying (255)), description (text), source_url (text), domain (character varying (50)), language (character varying (10)), and version (character). The data includes various datasets such as 'cleaning_set_of_th', 'LST', 'th_investing', 'gdb_or_th', 'moneylabstory.com', 'mti-insure.com', 'chilwithmeblog.com', 'esportivida.com', 'scb.co.th', 'chulabhornchannel_cra_ac_th', 'wongnai', 'criminal', 'hajai.com', 'tripitaka.mbu', 'pittanuej.com', 'wealthsolution_co.th', 'tripitaka_siamrath', 'taiwantourism.org', 'BEST Encyclopedia', 'tidlor.com', 'tcliv2', and 'govinigo.com'.

Figure 4.6 Dataset tables

4.4.2 Documents

The documents table stores the main textual data ingested from various sources. Each record represents a processed document with attributes such as dataset reference, content, language and metadata. The system successfully stores large volumes of text data extracted from multiple sources such as CC100, OSCAR and OpenSubtitles. The structured and cleaned text confirms that the ingestion and preprocessing pipeline successfully transforms raw data into AI-ready format.

The screenshot shows the PGAdmin interface with the 'public.documents' table selected. The table contains 22 rows of data. The columns are: id, dataset_id, parent_id, title (character varying (255)), and content (text). The data includes various documents such as 'filter_d10.filtered_0749734...', 'cc100_0518888.txt', 'cc100_10450579.txt', '825435', '2952608', 'cc100_3643432.txt', 'cc100_2778327.txt', '4189362', 'filter_d00.filtered_0950741...', 'filter_d02.filtered_0722277...', '3733402', '916500', 'cc100_12619669.txt', '6110456', 'cc100_4011708.txt', '2276958', 'filter_d02.filtered_0205149...', 'filter_d03.filtered_1253444...', '1874691', 'filter_d20.filtered_1322953...', and '4774088'.

Figure 4.7 Document Table

4.4.3 Labels

[illegible]

Figure 4.9 Labels Table

4.4.4 preprocessing_logs

The `preprocessing_logs` table records all steps performed during the data cleaning pipeline. This includes operations such as input validation, HTML removal and text normalization. Each log entry contains a step describing the preprocessing stage and a status indicating success or failure. Most preprocessing steps were successfully executed, with detailed logs recorded for traceability. This provides transparency and makes it easier to debug errors or analyze pipeline performance.

id	document_id	step	status	details
1	0000000f-c52e-4818-8188-51a8b91f19...	finalize_cleaned	success	['char_count': 1573, 'thai_ratio': 0.8614113159567705, 'reason': ...]
2	00000037-9aaa-49e4-bd6c-c542681e0...	validate_input	success	['original_chars': 1246]
3	0000005e-ab97-46f3-9482-49d83dec...	validate_input	success	['original_chars': 1053]
4	00000060-1d6c-4e58-a3f4-17e381db2...	remove_html	success	applied
5	00000067-a239-4339-9584-508452a2...	remove_html	success	applied
6	000000b6-1472-4111-b140-d61007b30...	finalize_cleaned	success	['char_count': 642, 'thai_ratio': 0.8286604361370716, 'reason': ...]
7	000000f1-81a5-48aa-8600-a7a112f05c...	validate_input	success	['original_chars': 1421]
8	00000127-e9f8-42be-9d69-cccae190c3...	remove_html	success	applied
9	00000164-408d-4c7e-9925-23b7e6dc...	validate_input	success	['original_chars': 1169]
10	000001da-f0b8-4715-aa56-2f982582d...	normalize_whitespace	success	no_change
11	000001f6-88cb-4fc2-a586-ec5a77f63f91	validate_input	success	['original_chars': 81]
12	000001f6-88cb-4fc2-a586-ec5a77f63f91	validate_input	success	['original_chars': 10842]
13	000001f6-88cb-4fc2-a586-ec5a77f63f91	validate_input	success	['original_chars': 3138]
14	00000218-7602-42eb-b6a8-576eb355e...	validate	success	
15	00000232-0868-4695-acae-8d8884e86...	finalize_cleaned	success	['char_count': 4377, 'thai_ratio': 0.912268677176148, 'reason': ...]
16	00000263-95df-4961-9bc6-8b993f8ae...	validate	success	
17	0000027e-fe2c-47ef-88db-401c2f1682...	remove_html	success	applied
18	000002b6-eece-4c75-a442-e043c26ac...	validate_input	success	['original_chars': 1223]
19	000002fe-81e8-4953-8d66-72f484e95b...	finalize_cleaned	success	['char_count': 285, 'thai_ratio': 0.7578947368421953, 'reason': ...]
20	00000305-036b-4f11-a723-c4d01d845...	normalize_whitespace	success	no_change
21	00000325-cf26-4778-9569-a53a978f5d...	remove_html	success	applied
22	00000340-5433-4b0c-8909-c9fec28a...	finalize_cleaned	success	['char_count': 47, 'thai_ratio': 0.6170212765957447, 'reason': ...]

Figure 4.10 Preprocessing_logs Table

4.5 Dashboard

The monitoring dashboard developed in this project plays a critical role in ensuring the reliability, performance and data quality of the entire data pipeline. By integrating multiple components including Airflow, MinIO and PostgreSQL, the system provides a comprehensive view of pipeline execution, storage usage and database performance.

4.5.1 Airflow Dashboards

For the Airflow cluster, the dashboard tracks DAG success/failure rate, task duration, number of running tasks and scheduler health. These panels help detect failed runs or unusually slow tasks.

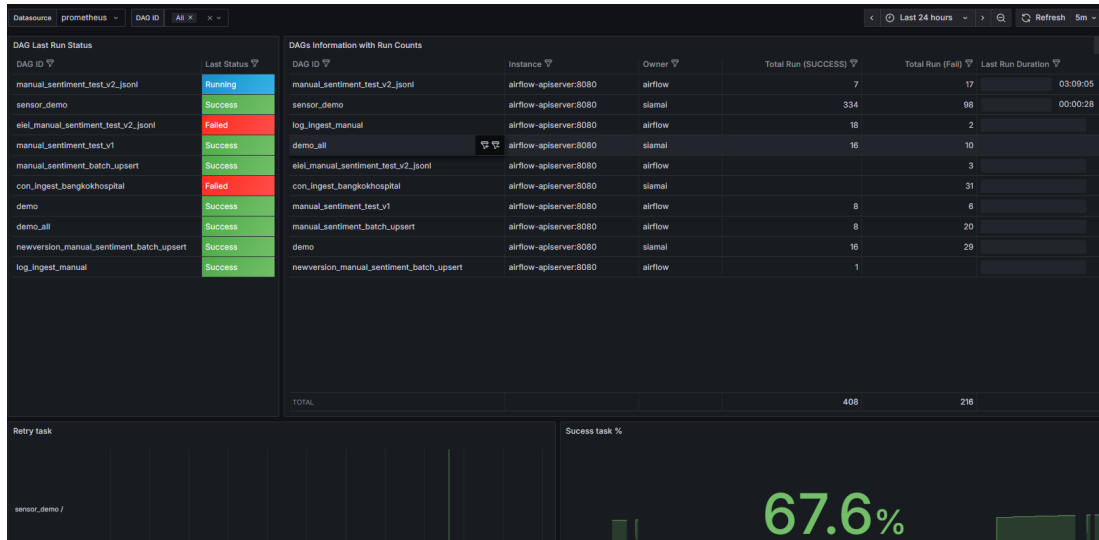


Figure 4.11 Airflow Dashboards

4.5.2 MinIO Dashboards

The MinIO cluster and bucket dashboards show bucket size, number of objects, upload/download throughput and storage usage. This view is used to check whether data ingestion is growing as expected and whether there is enough storage left.




Figure 4.12 MinIO Dashboards

4.5.3 PostgreSQL

The PostgreSQL database dashboard tracks query execution time, active connections, table size growth and other database metrics that may affect pipeline performance.




Figure 4.13 PostgreSQL Dashboards

4.5.4 Data Quality

This is a custom dashboard that reads metrics generated from SQL queries through the Postgres exporter. It shows batch counts, number of cleaned records and other indicators derived from the documents and preprocessing_log tables. These panels are used to check whether each batch was processed completely and to spot unusual drops or spikes in record counts.

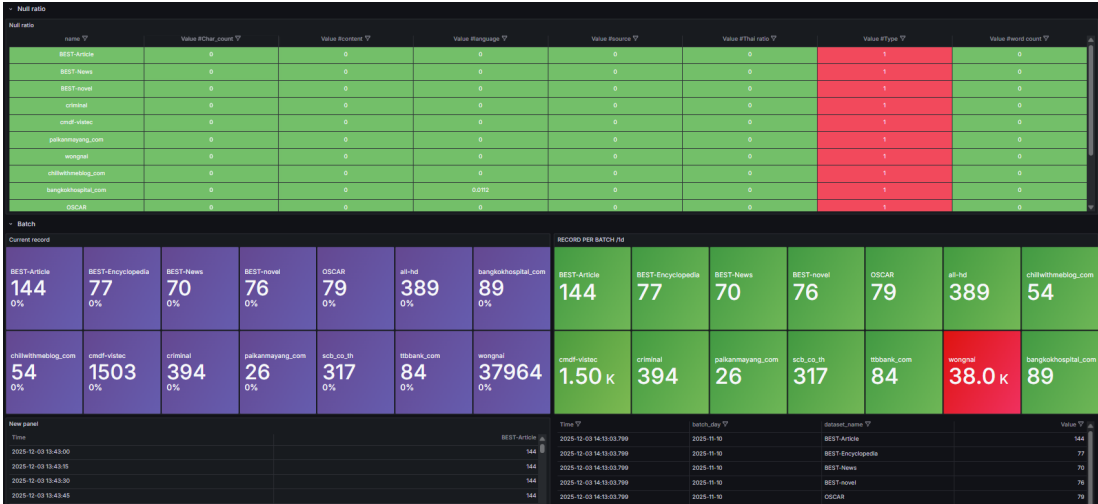


Figure 4.14 Data Quality Dashboards

CHAPTER 5 CONCLUSION

5.1 Conclusion

This project developed an Automated Data Ingestion and Preparation Pipeline for AI Model Training. The main goal of the project was to design and implement a pipeline that can collect, process, clean, validate, store and monitor different types of data so that the final outputs are suitable for downstream AI model training.

The completed system consists of several main components: a Data Lake using MinIO, automated processing pipelines using Apache Airflow, a PostgreSQL database for metadata and processed data storage, an unstructured data pipeline for PDF and audio processing, and a monitoring dashboard for system and data quality tracking. The MinIO data lake was organized into multiple layers to support raw, cleaned, processed and validated data storage. This structure helps separate each stage of the pipeline and makes the system easier to extend when new datasets are added.

For the unstructured pipeline, the PDF pipeline successfully processed 133 PDF documents with 7,346 pages, producing approximately 8.42 million characters of extracted Thai text. The audio pipeline processed 335,674 audio records, kept 307,793 valid transcripts and generated approximately 19.42 million characters from 656.93 hours of audio. In total, the PDF and audio pipelines produced approximately 27.85 million characters of AI-ready text data.

The PDF pipeline achieved a 100% document-level success rate and a 95.2% page-level success rate. The hybrid extraction method, which combines PyMuPDF and OCR, was effective because it allowed the system to process text-native PDFs efficiently while still extracting text from image-heavy PDFs. The audio pipeline also achieved strong results, with a 91.69% acceptance rate, a mean CER of 0.0447 and no failed or server-empty records in the final output. In addition, the audio format compatibility test confirmed that the pipeline could process 13 common audio formats, with 26 out of 26 format-clip trials successfully decoded and transcribed.

The system also includes a PostgreSQL database for storing datasets, documents, labels and preprocessing logs. This supports traceability, querying and future AI model training workflows. The dashboard provides visibility into Airflow, MinIO, PostgreSQL and data quality metrics. Overall, the project successfully achieved its main objective of building an automated pipeline that can transform raw data into clean, validated and AI-ready formats.

Table 5.1 Project completion status

Component	Description	Status
Data Lake	MinIO bucket structure for raw, cleaned, processed and validated data	Complete
PDF Pipeline	Extracts Thai text using PyMuPDF and OCR	Complete
Audio Pipeline	Converts Thai speech into text using ASR and filters by CER	Complete
Audio Format Compatibility	Supports WAV, MP3, M4A, OGG, WebM, WMA, 3GP, etc.	Complete
JSONL Pipeline	Processes JSONL/text datasets and stores cleaned outputs	Complete
Sentiment Processing	Adds sentiment labels to processed text data	Partially Complete
PostgreSQL Database	Stores datasets, documents, labels and preprocessing logs	Complete
Monitoring Dashboard	Tracks Airflow, MinIO, PostgreSQL and data quality metrics	Complete
Real-time Streaming Pipeline	Kafka-based real-time ingestion	Out of Scope
Model Training	Direct AI/LLM training and fine-tuning	Out of Scope

5.2 Problems, Obstacles and Solutions

During the project development, several problems and obstacles were encountered. These problems were mainly related to inconsistent data formats, Thai-language data quality and system performance when processing large-scale datasets. This section explains the main issues and how they were solved.

5.2.1 Inconsistent Data Format Across Sources

One major problem was that the input data came from many different formats and sources. The NECTEC text datasets included JSONL, TXT, CSV and PDF files. Some PDF documents were text-native, while others were scanned or image-heavy. The audio data also appeared in different container and codec formats such as WAV, MP3, M4A, OGG, WebM, WMA and 3GP. This inconsistency made it difficult to use one single processing method for all data.

To solve this problem, the system used a format-aware ingestion strategy. For PDFs, the pipeline first attempted direct text extraction using PyMuPDF. If the page did not contain usable embedded text, the pipeline used OCR as a fallback method. For audio data, the pipeline decoded different file formats using FFmpeg-backed tools and converted them into a common 16 kHz mono PCM WAV format before ASR processing. As a result, the ASR server only needed to process one standardized audio format.

The format compatibility test confirmed that this solution worked successfully. The pipeline processed 13 different audio formats across two test clips, and all 26 trials were successfully decoded and transcribed.

5.2.2 Data Quality Issues in Thai Language Processing

Another major challenge was data quality, especially for Thai-language processing. Thai text can be difficult to process because words are not always separated by spaces. In addition, OCR and ASR outputs may contain incorrect characters, repeated characters, missing tone marks, wrong transcriptions or hallucinated text.

For the PDF pipeline, some pages produced low-quality OCR outputs or no usable text at all. To prevent this, the system applied quality checks such as Thai-character ratio filtering, detection of abnormal repeated characters and filtering of corrupted or invalid characters. Empty pages and low-quality pages were excluded from the final output.

For the audio pipeline, transcription quality was evaluated using Character Error Rate (CER) instead of Word Error Rate (WER). CER was selected because it is more suitable for Thai text, where word boundaries are not always clear. Records with CER less than or equal to 0.20 were accepted. Silent audio detection was also added to remove clips that did not contain meaningful speech. These quality control methods helped improve the reliability of the final dataset.

5.2.3 System Performance and Scalability Issues

Processing large datasets created performance and scalability challenges. The PDF pipeline processed thousands of pages and OCR was computationally expensive. The audio pipeline processed more than 335,000 audio records, which required ASR inference and quality validation.

To reduce unnecessary computation, the PDF pipeline used PyMuPDF as the first extraction method and only used OCR when needed. For the audio pipeline, the system processed the dataset in Parquet shards and applied batch processing through the ASR workflow. The results across all 16 shards were consistent, with less than 1% variation in major metrics.

Monitoring was also added to help manage performance issues. The dashboard provides visibility into Airflow task status, MinIO storage usage, PostgreSQL performance and data quality metrics. This makes it easier to identify failed tasks, slow processing steps, storage growth and unusual changes in processed record counts.

5.3 Future Work and Suggestions

Although the project achieved its main objectives, there are several ways the system could be improved and extended in the future.

- **Vision-based PDF fallback:** The PDF pipeline could be improved by adding a stronger vision-based fallback method for pages that are currently classified as empty. A more advanced layout-aware OCR or document understanding model could help recover more content.
- **Audio quality tiers:** The audio pipeline could be extended by adding multiple quality tiers. Currently, the system uses a strict CER threshold of 0.20. Records with CER between 0.20 and 0.30 may still be usable for certain AI tasks, but they are excluded to keep the dataset precise.
- **CER and WER reporting:** The system could include both CER and WER reporting after integrating a reliable Thai word segmentation tool, especially for ASR benchmarking.
- **Unified database schema:** The schema could be further unified so that outputs from different modalities, such as PDF documents, JSONL records and audio transcripts, share a common structure with fields such as filename, source, content, quality score, processing method and validation status.
- **Detailed alerts and notifications:** The monitoring dashboard could be extended with more detailed alerts and automated notifications. The system could send alerts when a pipeline fails, when storage usage becomes too high, when the number of low-quality records increases suddenly or when processing time becomes abnormal.
- **Direct model training integration:** Future work could include direct integration with AI model training or fine-tuning workflows. The cleaned datasets could be connected to training frameworks such as Hugging Face Datasets, PyTorch or TensorFlow to support end-to-end AI development.

REFERENCES

1. Huyen, C. (2024). *AI engineering: Building applications with foundation models*. O'Reilly Media.
2. Airbyte. (n.d.). *Open-source data integration*. Retrieved August 20, 2025, from <https://airbyte.com/>
3. Amazon Web Services. (n.d.). *Amazon S3*. Retrieved August 22, 2025, from <https://aws.amazon.com/s3/>
4. Apache Airflow. (n.d.). *Apache Airflow documentation*. Retrieved August 22, 2025, from <https://airflow.apache.org/docs/apache-airflow/stable/>
5. Apache Arrow. (n.d.). *Apache Arrow documentation*. Retrieved August 18, 2025, from <https://arrow.apache.org/>
6. Apache Kafka. (n.d.). *Apache Kafka documentation*. Retrieved August 18, 2025, from <https://kafka.apache.org/>
7. Apache NiFi. (n.d.). *Apache NiFi project*. Retrieved August 18, 2025, from <https://nifi.apache.org/>
8. Chacon, S., & Straub, B. (2014). *Pro Git* (2nd ed.). Apress. <https://git-scm.com/book/en/v2>
9. Docker. (n.d.). *What is Docker?* Retrieved August 13, 2025, from <https://www.docker.com/resources/what-container>
10. Fortinet. (n.d.). *FortiClient VPN*. Retrieved August 14, 2025, from <https://www.fortinet.com/products/endpoint-security/forticlient>
11. Grafana Labs. (n.d.). *What is Grafana?* Retrieved August 19, 2025, from <https://grafana.com/>
12. Google Cloud. (n.d.). *BigQuery*. Retrieved August 15, 2025, from <https://cloud.google.com/bigquery>
13. JaiedAI. (n.d.). *EasyOCR*. Retrieved September 17, 2025, from <https://github.com/JaiedAI/EasyOCR>
14. McKinney, W. (2010). Data structures for statistical computing in Python. *Proceedings of the 9th Python in Science Conference*, 56–61. <https://pandas.pydata.org/>
15. Microsoft. (n.d.). *Visual Studio Code*. Retrieved August 13, 2025, from <https://code.visualstudio.com/>
16. MinIO. (n.d.). *High performance object storage*. Retrieved August 14, 2025, from <https://min.io/>
17. PaddleOCR Contributors. (n.d.). *PaddleOCR*. Retrieved September 12, 2025, from <https://github.com/PaddlePaddle/PaddleOCR>
18. PostgreSQL Global Development Group. (n.d.). *PostgreSQL: The world's most advanced open source database*. Retrieved August 18, 2025, from <https://www.postgresql.org/>
19. Prometheus Authors. (n.d.). *Prometheus monitoring system*. Retrieved August 18, 2025, from <https://prometheus.io/>
20. Phatthiyaphaibun, W., et al. (2023). PyThaiNLP: Thai natural language processing in Python. *Proceedings of the 3rd Workshop for NLP-OSS*, 25–36. <https://aclanthology.org/2023.nlposs-1.4/>
21. Smith, R. (2007). An overview of the Tesseract OCR engine. *Proceedings of the International Conference on Document Analysis and Recognition (ICDAR)*, 629–633. <https://doi.org/10.1109/ICDAR.2007.4376991>
22. Typhoon OCR. (n.d.). *Typhoon OCR project*. Retrieved September 15, 2025, from <https://github.com/scb-10x/typhoon-ocr>
23. Typhoon AI. (n.d.). *Typhoon Isan ASR Whisper*. Hugging Face. Retrieved October 10, 2025, from <https://huggingface.co/typhoon-ai/typhoon-isan-asr-whisper>

24. Typhoon AI. (n.d.). *Typhoon OCR 7B*. Hugging Face. Retrieved September 15, 2025, from <https://huggingface.co/typhoon-ai/typhoon-ocr-7b>
25. NVIDIA Corporation. (n.d.). *NVIDIA DGX B200*. Retrieved September 15, 2025, from <https://www.nvidia.com/en-us/data-center/dgx-b200/>
26. AIRsearch. (2021, March 3). *WangchanBERTa: Pre-trained Thai language model*. Retrieved October 4, 2025, from <https://airesearch.in.th/releases/wangchanberta-pre-trained-thai-language-model/>